

Week 5: Bayesian linear regression and introduction to Stan

11/02/23

Introduction

Today we will be starting off using Stan, looking at the kid's test score data set (available in resources for the [Gelman Hill textbook](#)).

```
library(tidyverse)
library(rstan)
library(tidybayes)
library(here)
```

The data look like this:

```
kidiq <- read_rds(here("data","kidiq.RDS"))
kidiq
```

```
# A tibble: 434 x 4
  kid_score mom_hs mom_iq mom_age
  <int>    <dbl>  <dbl>   <int>
1      65      1  121.     27
2      98      1   89.4     25
3      85      1  115.     27
4      83      1   99.4     25
5     115      1   92.7     27
6      98      0  108.     18
7      69      1  139.     20
8     106      1  125.     23
9     102      1   81.6     24
```

```
10      95      1   95.1      19
# ... with 424 more rows
```

As well as the kid's test scores, we have a binary variable indicating whether or not the mother completed high school, the mother's IQ and age.

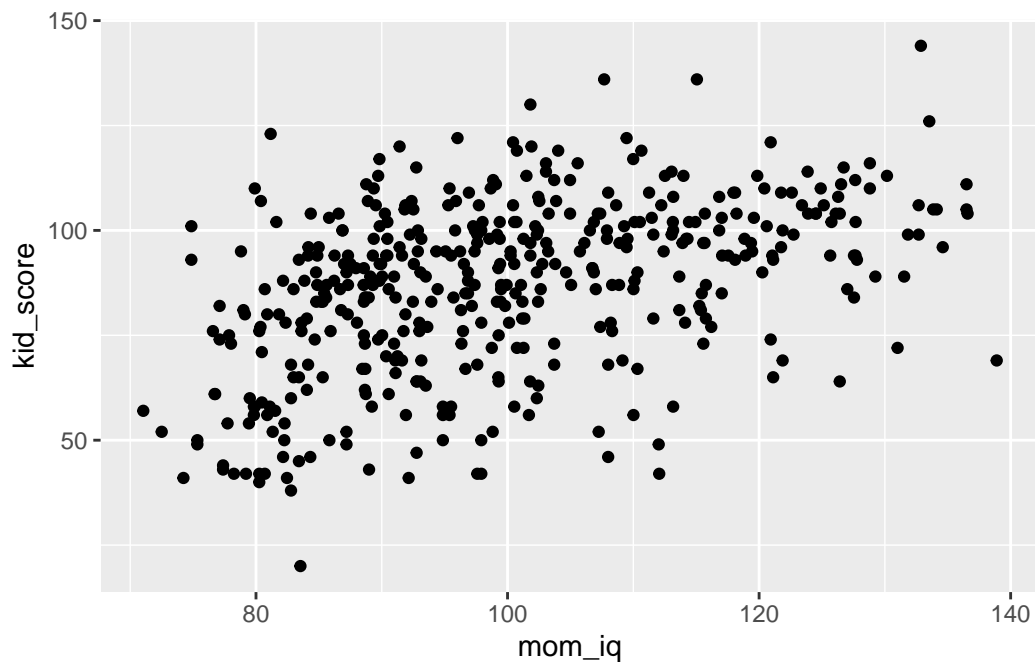
Descriptives

Question 1

Use plots or tables to show three interesting observations about the data. Remember:

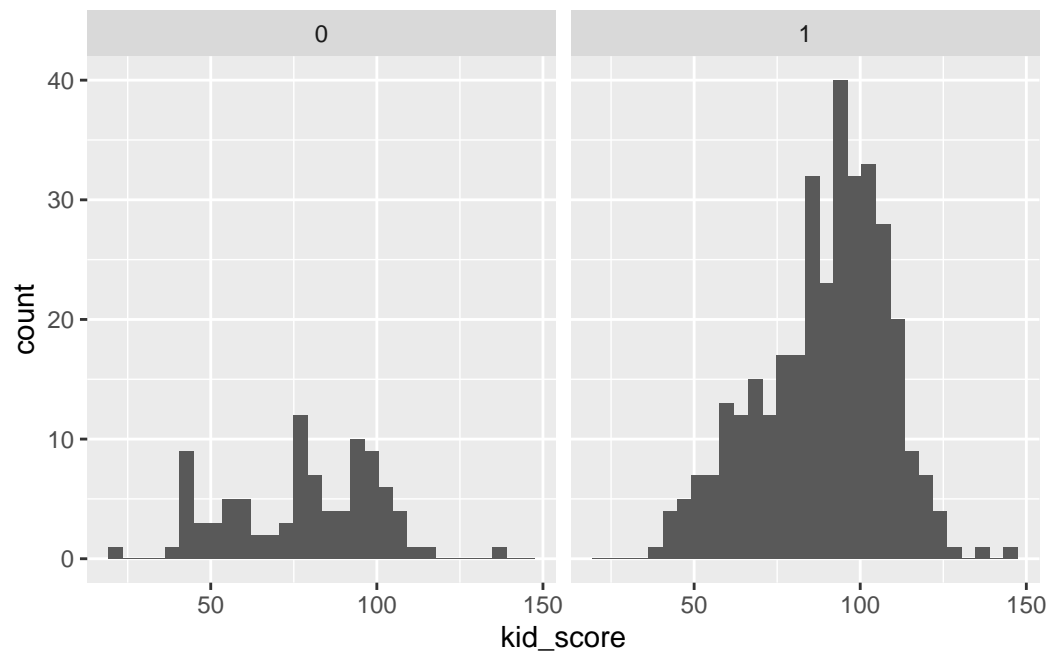
- Explain what your graph/ tables show
- Choose a graph type that's appropriate to the data type

```
ggplot(data=kidiq)+
  geom_point(aes(x=mom_iq, y=kid_score))
```




```
ggplot(data=kidiq)+
  geom_histogram(aes(kid_score))+
```

```
facet_grid(~mom_hs)
```



```
ggplot(data=kidiq)+  
  geom_point(aes(x=mom_age, y=kid_score))
```



1. Since both `kid_score` and `mom_iq` are numerical, I first plot a scatter plot of kid's IQ score vs mom's IQ score, and found that kid's IQ score increases as mom's IQ score increases.
2. Since `mom_hs` is binary and `kid_score` is numerical, I then plot histograms of kid's IQ by mom's high school status, and found that for mother with a high school degree, the kid's IQ has larger density on higher IQ scores than for mother without a high school degree.
3. Since both `kid_score` and `mom_age` are numerical, I also plot a scatter plot of kid's IQ score vs mom's age, and found that there is no clear pattern between kid's IQ and mom's age.

Estimating mean, no covariates

In class we were trying to estimate the mean and standard deviation of the kid's test scores. The `kids2.stan` file contains a Stan model to do this. If you look at it, you will notice the first `data` chunk lists some inputs that we have to define: the outcome variable `y`, number of observations `N`, and the mean and standard deviation of the prior on `mu`. Let's define all these values in a `data` list.

```

y <- kidiq$kid_score
mu0 <- 80
sigma0 <- 10

# named list to input for stan function
data <- list(y = y,
             N = length(y),
             mu0 = mu0,
             sigma0 = sigma0)

```

Now we can run the model:

```

fit <- stan(file = here("code/models/kids2.stan"),
            data = data,
            chains = 3,
            iter = 500,
            seed = 1)

```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 2.5e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.25 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 500 [0%] (Warmup)

Chain 1: Iteration: 50 / 500 [10%] (Warmup)

Chain 1: Iteration: 100 / 500 [20%] (Warmup)

Chain 1: Iteration: 150 / 500 [30%] (Warmup)

Chain 1: Iteration: 200 / 500 [40%] (Warmup)

Chain 1: Iteration: 250 / 500 [50%] (Warmup)

Chain 1: Iteration: 251 / 500 [50%] (Sampling)

Chain 1: Iteration: 300 / 500 [60%] (Sampling)

Chain 1: Iteration: 350 / 500 [70%] (Sampling)

Chain 1: Iteration: 400 / 500 [80%] (Sampling)

Chain 1: Iteration: 450 / 500 [90%] (Sampling)

Chain 1: Iteration: 500 / 500 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.010573 seconds (Warm-up)

Chain 1: 0.004751 seconds (Sampling)

Chain 1: 0.015324 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 6e-06 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.06 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 500 [0%] (Warmup)

Chain 2: Iteration: 50 / 500 [10%] (Warmup)

Chain 2: Iteration: 100 / 500 [20%] (Warmup)

Chain 2: Iteration: 150 / 500 [30%] (Warmup)

Chain 2: Iteration: 200 / 500 [40%] (Warmup)

Chain 2: Iteration: 250 / 500 [50%] (Warmup)

Chain 2: Iteration: 251 / 500 [50%] (Sampling)

Chain 2: Iteration: 300 / 500 [60%] (Sampling)

Chain 2: Iteration: 350 / 500 [70%] (Sampling)

Chain 2: Iteration: 400 / 500 [80%] (Sampling)

Chain 2: Iteration: 450 / 500 [90%] (Sampling)

Chain 2: Iteration: 500 / 500 [100%] (Sampling)

Chain 2:

Chain 2: Elapsed Time: 0.007538 seconds (Warm-up)

Chain 2: 0.005284 seconds (Sampling)

Chain 2: 0.012822 seconds (Total)

Chain 2:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 3).

Chain 3:

Chain 3: Gradient evaluation took 6e-06 seconds

Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.06 seconds.

Chain 3: Adjust your expectations accordingly!

Chain 3:

Chain 3:

Chain 3: Iteration: 1 / 500 [0%] (Warmup)

Chain 3: Iteration: 50 / 500 [10%] (Warmup)

Chain 3: Iteration: 100 / 500 [20%] (Warmup)

Chain 3: Iteration: 150 / 500 [30%] (Warmup)

Chain 3: Iteration: 200 / 500 [40%] (Warmup)

Chain 3: Iteration: 250 / 500 [50%] (Warmup)

Chain 3: Iteration: 251 / 500 [50%] (Sampling)

Chain 3: Iteration: 300 / 500 [60%] (Sampling)

Chain 3: Iteration: 350 / 500 [70%] (Sampling)

```
Chain 3: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 3: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 3: Iteration: 500 / 500 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.007261 seconds (Warm-up)
Chain 3:           0.005279 seconds (Sampling)
Chain 3:           0.01254 seconds (Total)
Chain 3:
```

Look at the summary

```
fit
```

Inference for Stan model: kids2.

3 chains, each with iter=500; warmup=250; thin=1;

post-warmup draws per chain=250, total post-warmup draws=750.

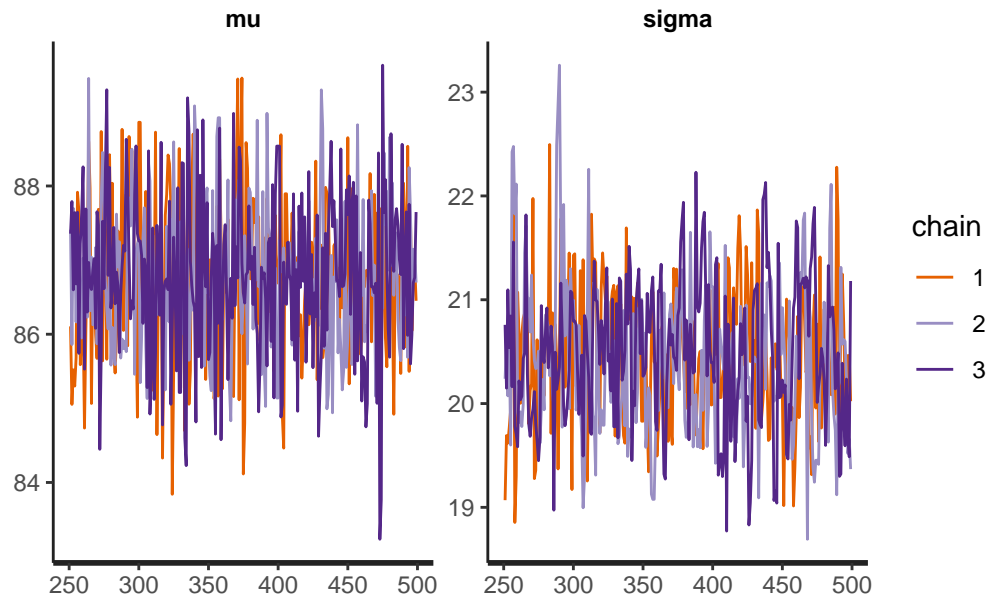
	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
mu	86.79	0.04	1.04	84.81	86.07	86.80	87.48	88.77	743
sigma	20.45	0.04	0.69	19.17	19.97	20.44	20.92	21.87	315
lp__	-1525.83	0.06	1.07	-1528.49	-1526.26	-1525.51	-1525.05	-1524.78	369
Rhat									
mu	1.00								
sigma	1.00								
lp__	1.02								

Samples were drawn using NUTS(diag_e) at Sat Feb 11 12:23:26 2023.

For each parameter, n_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

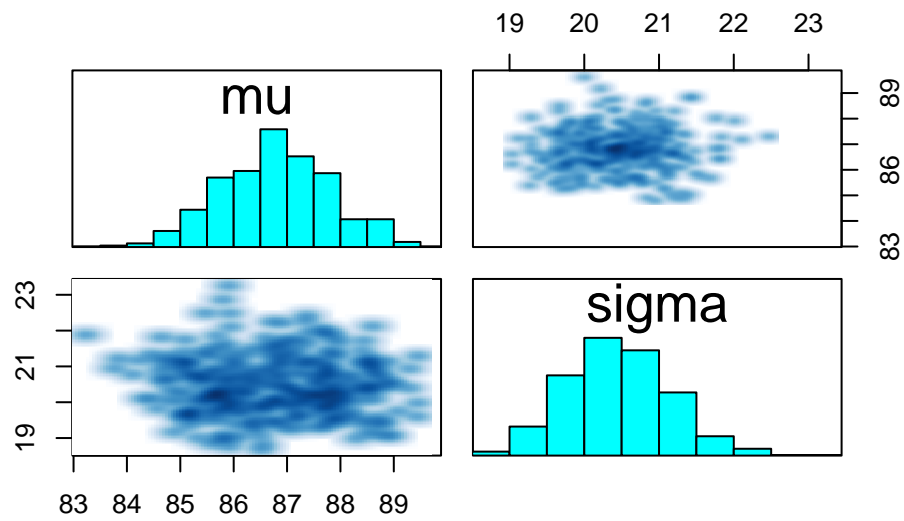
Traceplot

```
traceplot(fit)
```

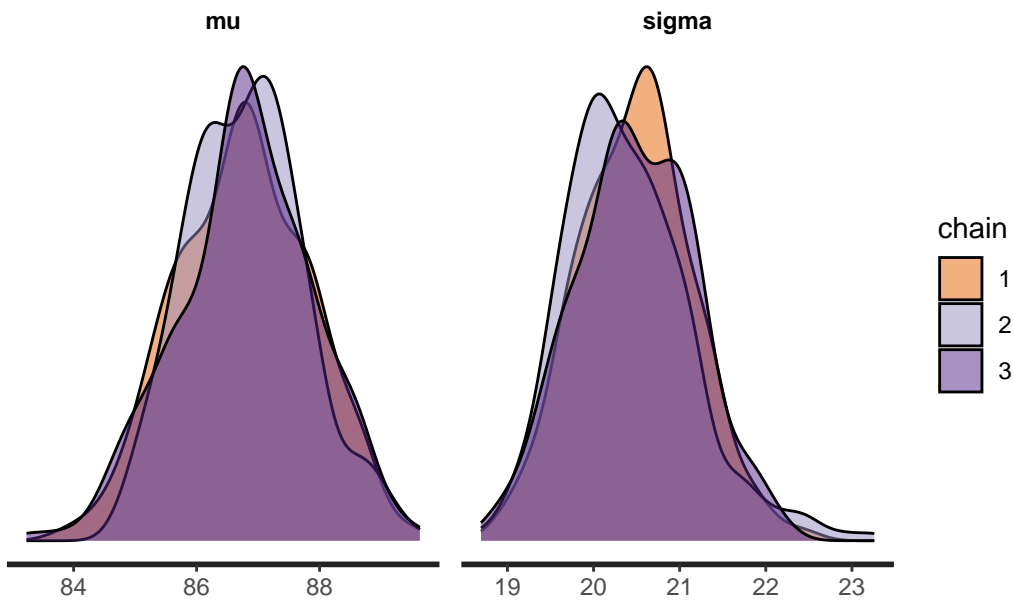


All looks fine.

```
pairs(fit, pars = c("mu", "sigma"))
```

```
stan_dens(fit, separate_chains = TRUE)
```



Understanding output

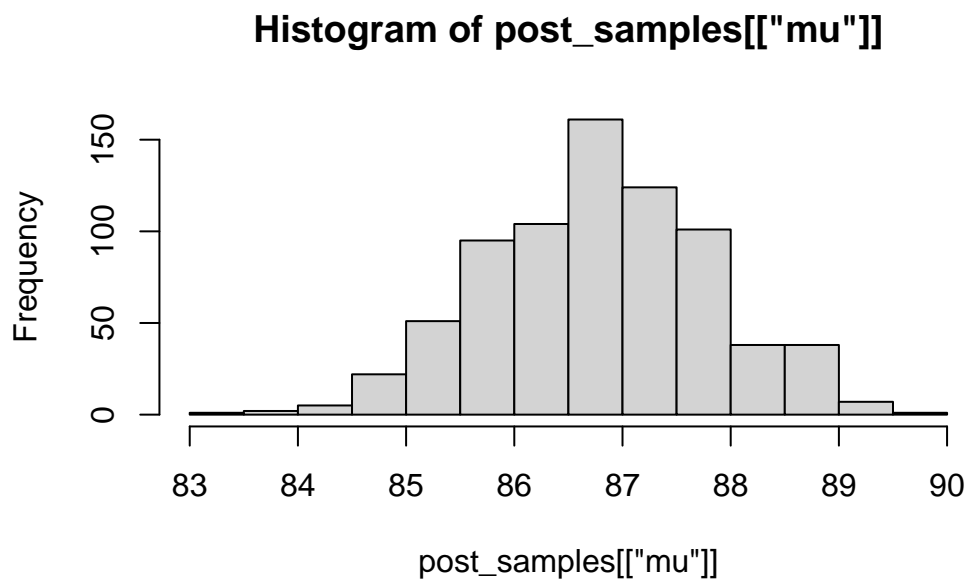
What does the model actually give us? A number of samples from the posteriors. To see this, we can use `extract` to get the samples.

```
post_samples <- extract(fit)
head(post_samples[["mu"]])
```

```
[1] 86.21779 85.47105 87.83292 85.85235 86.35743 87.39406
```

This is a list, and in this case, each element of the list has 4000 samples. E.g. quickly plot a histogram of `mu`

```
hist(post_samples[["mu"]])
```



```
median(post_samples[["mu"]])
```

```
[1] 86.79872
```

```
# 95% bayesian credible interval
quantile(post_samples[["mu"]], 0.025)
```

```
2.5%
84.81158
```

```
quantile(post_samples[["mu"]], 0.975)
```

```
97.5%
88.76504
```

Plot estimates

There are a bunch of packages, built-in functions that let you plot the estimates from the model, and I encourage you to explore these options (particularly in **bayesplot**, which we will most likely be using later on). I like using the **tidybayes** package, which allows us to easily get the posterior samples in a tidy format (e.g. using `gather_draws` to get in long format). Once we have that, it's easy to just pipe and do ggplots as usual.

Get the posterior samples for mu and sigma in long format:

```
dsamples <- fit |>
  gather_draws(mu, sigma) # gather = long format
dsamples
```

```
# A tibble: 1,500 x 5
# Groups:   .variable [2]
  .chain .iteration .draw .variable .value
  <int>     <int> <int> <chr>     <dbl>
1       1         1     1 mu         86.1
2       1         2     2 mu         85.1
3       1         3     3 mu         85.5
4       1         4     4 mu         85.3
5       1         5     5 mu         85.7
6       1         6     6 mu         87.9
7       1         7     7 mu         87.7
8       1         8     8 mu         85.6
9       1         9     9 mu         86.0
10      1        10    10 mu         86.2
# ... with 1,490 more rows
```

```
# wide format
fit |> spread_draws(mu, sigma)

# A tibble: 750 x 5
  .chain .iteration .draw    mu sigma
  <int>      <int> <int> <dbl> <dbl>
1       1         1     1  86.1  19.1
2       1         2     2  85.1  19.7
3       1         3     3  85.5  19.6
4       1         4     4  85.3  19.8
5       1         5     5  85.7  19.9
6       1         6     6  87.9  21.4
7       1         7     7  87.7  22.1
8       1         8     8  85.6  18.9
9       1         9     9  86.0  19.3
10      1        10    10  86.2  21.1
# ... with 740 more rows
```

```
# quickly calculate the quantiles using
```

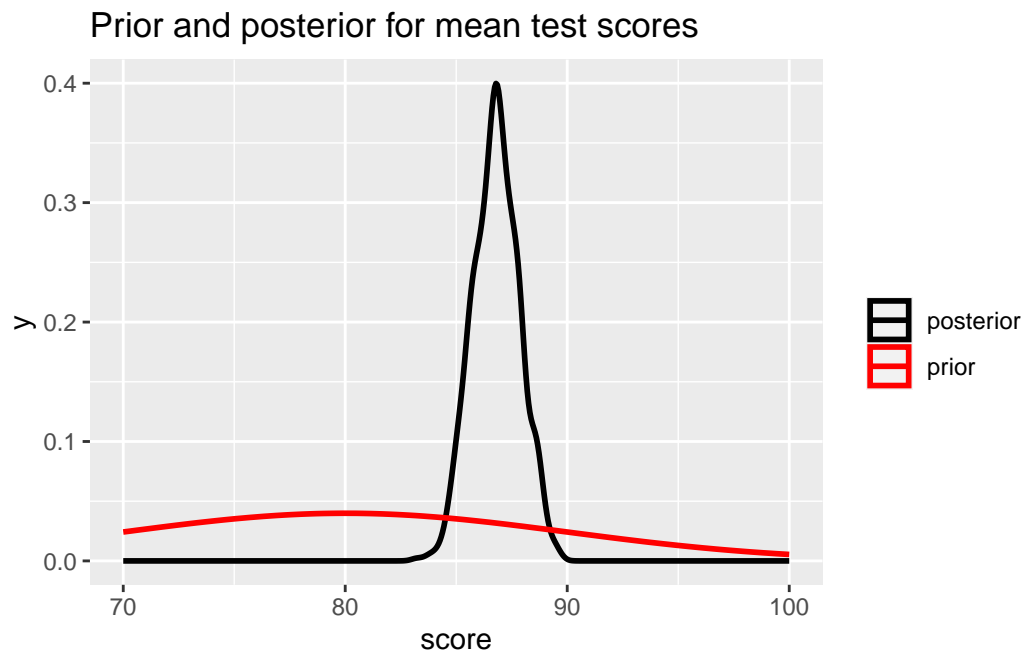
```
dsamples |>
  median_qi(.width = 0.8)
```

```
# A tibble: 2 x 7
  .variable .value .lower .upper .width .point .interval
  <chr>      <dbl> <dbl> <dbl> <dbl> <chr> <chr>
1 mu        86.8  85.5  88.1  0.8 median qi
2 sigma     20.4  19.6  21.3  0.8 median qi
```

Let's plot the density of the posterior samples for mu and add in the prior distribution

```
dsamples |>
  filter(.variable == "mu") |>
  ggplot(aes(.value, color = "posterior")) + geom_density(size = 1) +
  xlim(c(70, 100)) +
  stat_function(fun = dnorm,
    args = list(mean = mu0,
                 sd = sigma0),
    aes(colour = 'prior'), size = 1) +
  scale_color_manual(name = "", values = c("prior" = "red", "posterior" = "black")) +
```

```
ggtitle("Prior and posterior for mean test scores") +
xlab("score")
```



Question 2

Change the prior to be much more informative (by changing the standard deviation to be 0.1). Rerun the model. Do the estimates change? Plot the prior and posterior densities.

```
y <- kidiq$kid_score
mu0 <- 80
sigma0 <- 0.1

# named list to input for stan function
data <- list(y = y,
             N = length(y),
             mu0 = mu0,
             sigma0 = sigma0)

fit <- stan(file = here("code/models/kids2.stan"),
            data = data,
            chains = 3,
```

```
iter = 500,  
seed = 1)
```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 7e-06 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.07 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 500 [0%] (Warmup)

Chain 1: Iteration: 50 / 500 [10%] (Warmup)

Chain 1: Iteration: 100 / 500 [20%] (Warmup)

Chain 1: Iteration: 150 / 500 [30%] (Warmup)

Chain 1: Iteration: 200 / 500 [40%] (Warmup)

Chain 1: Iteration: 250 / 500 [50%] (Warmup)

Chain 1: Iteration: 251 / 500 [50%] (Sampling)

Chain 1: Iteration: 300 / 500 [60%] (Sampling)

Chain 1: Iteration: 350 / 500 [70%] (Sampling)

Chain 1: Iteration: 400 / 500 [80%] (Sampling)

Chain 1: Iteration: 450 / 500 [90%] (Sampling)

Chain 1: Iteration: 500 / 500 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.006578 seconds (Warm-up)

Chain 1: 0.005536 seconds (Sampling)

Chain 1: 0.012114 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 6e-06 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.06 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 500 [0%] (Warmup)

Chain 2: Iteration: 50 / 500 [10%] (Warmup)

Chain 2: Iteration: 100 / 500 [20%] (Warmup)

Chain 2: Iteration: 150 / 500 [30%] (Warmup)

Chain 2: Iteration: 200 / 500 [40%] (Warmup)

Chain 2: Iteration: 250 / 500 [50%] (Warmup)

```

Chain 2: Iteration: 251 / 500 [ 50%] (Sampling)
Chain 2: Iteration: 300 / 500 [ 60%] (Sampling)
Chain 2: Iteration: 350 / 500 [ 70%] (Sampling)
Chain 2: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 2: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 2: Iteration: 500 / 500 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.005947 seconds (Warm-up)
Chain 2:           0.004947 seconds (Sampling)
Chain 2:           0.010894 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 6e-06 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.06 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:   1 / 500 [  0%] (Warmup)
Chain 3: Iteration:  50 / 500 [ 10%] (Warmup)
Chain 3: Iteration: 100 / 500 [ 20%] (Warmup)
Chain 3: Iteration: 150 / 500 [ 30%] (Warmup)
Chain 3: Iteration: 200 / 500 [ 40%] (Warmup)
Chain 3: Iteration: 250 / 500 [ 50%] (Warmup)
Chain 3: Iteration: 251 / 500 [ 50%] (Sampling)
Chain 3: Iteration: 300 / 500 [ 60%] (Sampling)
Chain 3: Iteration: 350 / 500 [ 70%] (Sampling)
Chain 3: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 3: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 3: Iteration: 500 / 500 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.006061 seconds (Warm-up)
Chain 3:           0.005401 seconds (Sampling)
Chain 3:           0.011462 seconds (Total)
Chain 3:

```

`fit`

Inference for Stan model: kids2.

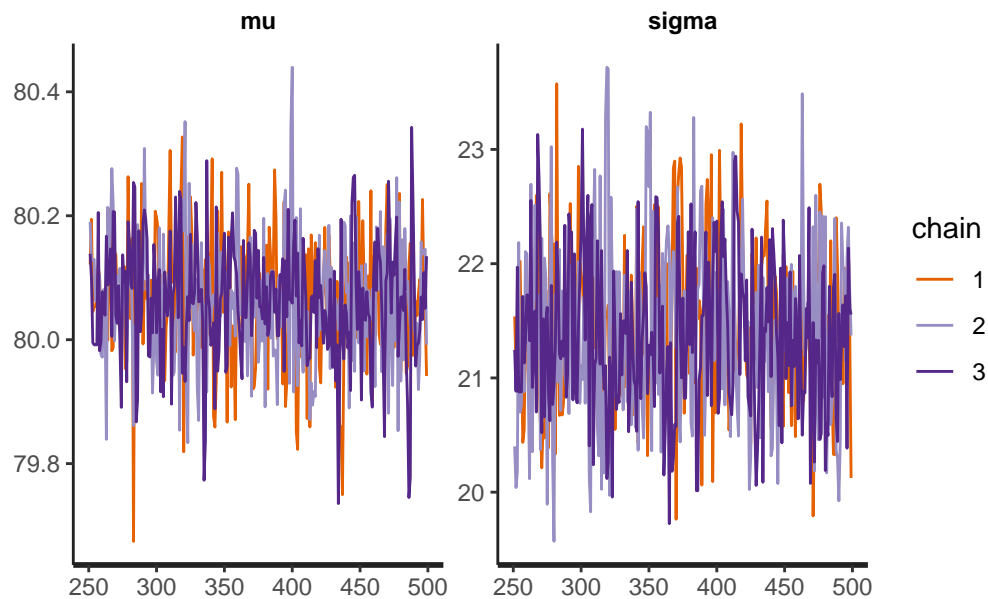
3 chains, each with iter=500; warmup=250; thin=1;
post-warmup draws per chain=250, total post-warmup draws=750.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
mu	80.06	0.00	0.10	79.86	80.00	80.07	80.12	80.25	669
sigma	21.42	0.03	0.72	20.09	20.90	21.40	21.90	22.90	618
lp__	-1548.35	0.05	1.03	-1551.09	-1548.67	-1548.04	-1547.67	-1547.40	352

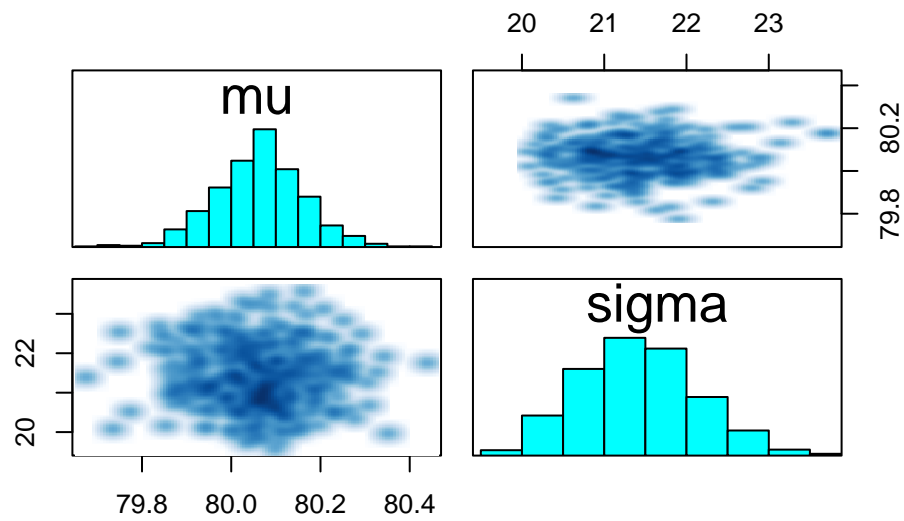
	Rhat
mu	1
sigma	1
lp__	1

Samples were drawn using NUTS(diag_e) at Sat Feb 11 12:23:28 2023.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

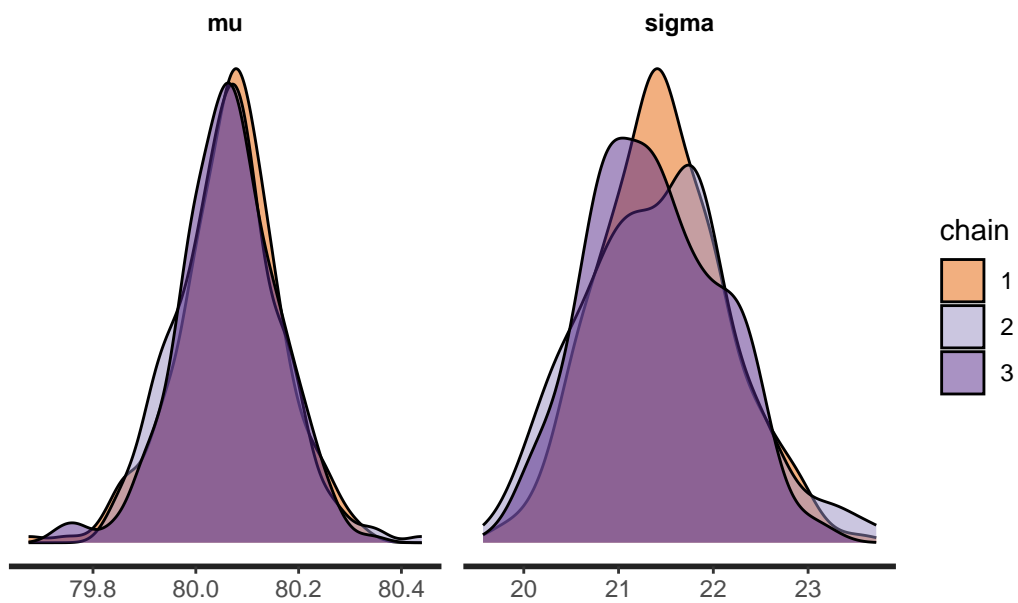
```
traceplot(fit)
```



```
pairs(fit, pars = c("mu", "sigma"))
```

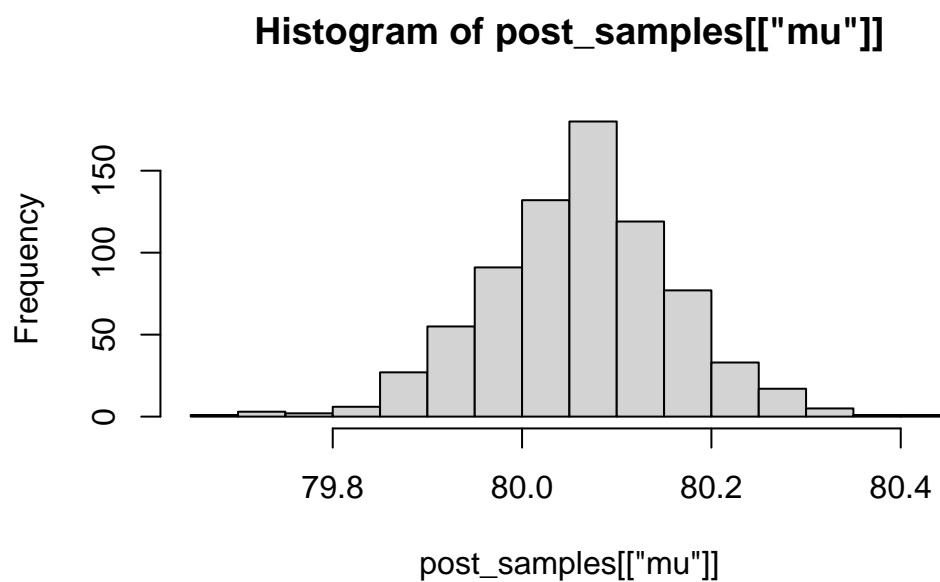
```
stan_dens(fit, separate_chains = TRUE)
```



```
post_samples <- extract(fit)
head(post_samples[["mu"]])
```

```
[1] 80.23602 80.01066 80.10293 80.00399 80.09156 80.05373
```

```
hist(post_samples[["mu"]])
```



```
median(post_samples[["mu"]])
```

```
[1] 80.06706
```

```
# 95% bayesian credible interval
quantile(post_samples[["mu"]], 0.025)
```

```
2.5%
79.86062
```

```
quantile(post_samples[["mu"]], 0.975)
```

```
97.5%  
80.25431
```

```
dsamples <- fit |>  
  gather_draws(mu, sigma) # gather = long format  
dsamples
```

```
# A tibble: 1,500 x 5  
# Groups:   .variable [2]  
  .chain .iteration .draw .variable .value  
  <int>      <int> <int> <chr>      <dbl>  
1      1         1     1 mu         80.1  
2      1         2     2 mu         80.2  
3      1         3     3 mu         80.0  
4      1         4     4 mu         80.1  
5      1         5     5 mu         80.1  
6      1         6     6 mu         80.1  
7      1         7     7 mu         80.1  
8      1         8     8 mu         80.1  
9      1         9     9 mu         80.0  
10     1        10    10 mu         80.1  
# ... with 1,490 more rows
```

```
# wide format  
fit |> spread_draws(mu, sigma)
```

```
# A tibble: 750 x 5  
  .chain .iteration .draw    mu sigma  
  <int>      <int> <int> <dbl> <dbl>  
1      1         1     1  80.1  21.5  
2      1         2     2  80.2  21.1  
3      1         3     3  80.0  21.0  
4      1         4     4  80.1  22.0  
5      1         5     5  80.1  22.0  
6      1         6     6  80.1  21.3  
7      1         7     7  80.1  20.4  
8      1         8     8  80.1  20.5
```

```

  9      1      9      9 80.0 21.4
10      1     10     10 80.1 21.7
# ... with 740 more rows

```

```
# quickly calculate the quantiles using
```

```

dsamples |>
  median_qi(.width = 0.8)

```

```

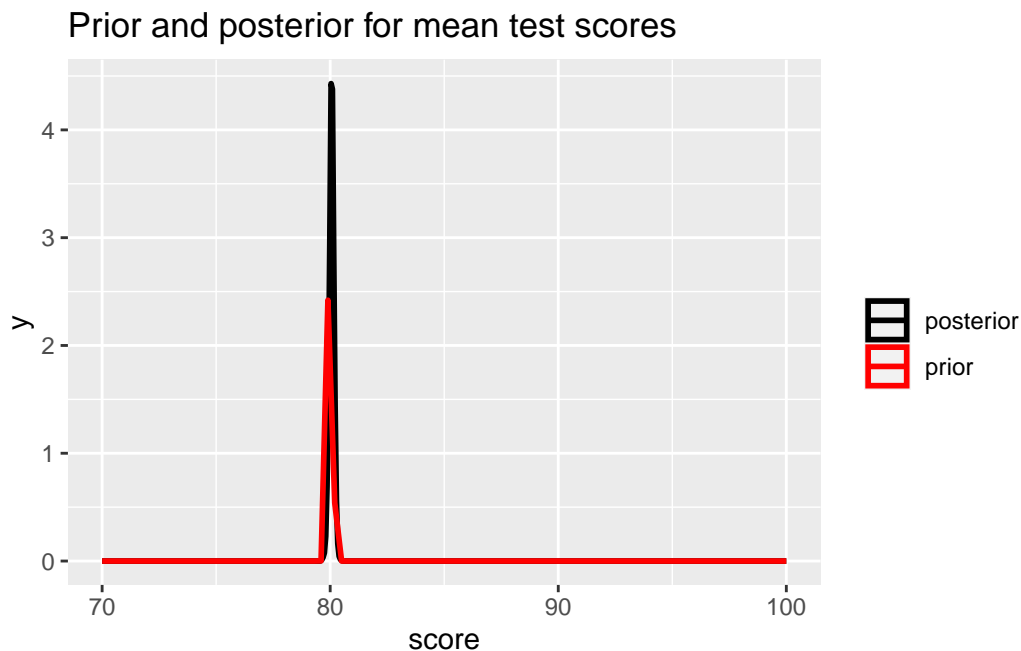
# A tibble: 2 x 7
  .variable .value .lower .upper .width .point .interval
  <chr>      <dbl> <dbl> <dbl> <dbl> <chr> <chr>
1 mu        80.1  79.9  80.2   0.8 median qi
2 sigma     21.4  20.5  22.4   0.8 median qi

```

```

dsamples |>
  filter(.variable == "mu") |>
  ggplot(aes(.value, color = "posterior")) + geom_density(size = 1) +
  xlim(c(70, 100)) +
  stat_function(fun = dnorm,
    args = list(mean = mu0,
                sd = sigma0),
    aes(colour = 'prior'), size = 1) +
  scale_color_manual(name = "", values = c("prior" = "red", "posterior" = "black")) +
  ggtitle("Prior and posterior for mean test scores") +
  xlab("score")

```



The estimates changed. Previously, the median of μ is 86.80 and median of σ is 20.44, but now the median of μ is 80.07 and median of σ is 21.40.

Adding covariates

Now let's see how kid's test scores are related to mother's education. We want to run the simple linear regression

$$Score = \alpha + \beta X$$

where $X = 1$ if the mother finished high school and zero otherwise.

`kid3.stan` has the stan model to do this. Notice now we have some inputs related to the design matrix X and the number of covariates (in this case, it's just 1).

Let's get the data we need and run the model.

```
X <- as.matrix(kidiq$mom_hs, ncol = 1) # force this to be a matrix
K <- 1

data <- list(y = y, N = length(y),
             X = X, K = K)
```

```
fit2 <- stan(file = here("code/models/kids3.stan"),
             data = data,
             iter = 1000,
             seed = 1)
```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 5.5e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.55 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 1000 [0%] (Warmup)

Chain 1: Iteration: 100 / 1000 [10%] (Warmup)

Chain 1: Iteration: 200 / 1000 [20%] (Warmup)

Chain 1: Iteration: 300 / 1000 [30%] (Warmup)

Chain 1: Iteration: 400 / 1000 [40%] (Warmup)

Chain 1: Iteration: 500 / 1000 [50%] (Warmup)

Chain 1: Iteration: 501 / 1000 [50%] (Sampling)

Chain 1: Iteration: 600 / 1000 [60%] (Sampling)

Chain 1: Iteration: 700 / 1000 [70%] (Sampling)

Chain 1: Iteration: 800 / 1000 [80%] (Sampling)

Chain 1: Iteration: 900 / 1000 [90%] (Sampling)

Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.12928 seconds (Warm-up)

Chain 1: 0.096947 seconds (Sampling)

Chain 1: 0.226227 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 2.3e-05 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.23 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 1000 [0%] (Warmup)

Chain 2: Iteration: 100 / 1000 [10%] (Warmup)

Chain 2: Iteration: 200 / 1000 [20%] (Warmup)

Chain 2: Iteration: 300 / 1000 [30%] (Warmup)

```

Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.136823 seconds (Warm-up)
Chain 2:                0.085941 seconds (Sampling)
Chain 2:                0.222764 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 2.4e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.24 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:   1 / 1000 [  0%] (Warmup)
Chain 3: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 3: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 3: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 3: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 3: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 3: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 3: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 3: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 3: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 3: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 3: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.167209 seconds (Warm-up)
Chain 3:                0.096508 seconds (Sampling)
Chain 3:                0.263717 seconds (Total)
Chain 3:

```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 4).

```

Chain 4:
Chain 4: Gradient evaluation took 2.4e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.24 seconds.

```

```

Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration: 1 / 1000 [ 0%] (Warmup)
Chain 4: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 4: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 4: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 4: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 4: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 4: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 4: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 4: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 4: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 4: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 4: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 0.124515 seconds (Warm-up)
Chain 4: 0.08831 seconds (Sampling)
Chain 4: 0.212825 seconds (Total)
Chain 4:

```

Question 3

- a) Confirm that the estimates of the intercept and slope are comparable to results from `lm()`

```
fit2
```

Inference for Stan model: kids3.

4 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=2000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
alpha	78.00	0.07	1.94	74.12	76.74	78.06	79.33	81.67
beta[1]	11.18	0.08	2.19	7.00	9.66	11.17	12.60	15.56
sigma	19.83	0.02	0.66	18.59	19.39	19.81	20.25	21.23
lp__	-1514.33	0.04	1.21	-1517.35	-1514.93	-1514.03	-1513.44	-1512.98
	n_eff	Rhat						
alpha	804	1.01						
beta[1]	846	1.01						
sigma	1042	1.00						


```
lp__      758 1.00
```

Samples were drawn using NUTS(diag_e) at Sat Feb 11 12:23:53 2023.
For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

```
model<-lm(kid_score~as.factor(mom_hs), data=kidiq)
summary(model)
```

Call:

```
lm(formula = kid_score ~ as.factor(mom_hs), data = kidiq)
```

Residuals:

Min	1Q	Median	3Q	Max
-57.55	-13.32	2.68	14.68	58.45

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	77.548	2.059	37.670	< 2e-16 ***
as.factor(mom_hs)1	11.771	2.322	5.069	5.96e-07 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.85 on 432 degrees of freedom

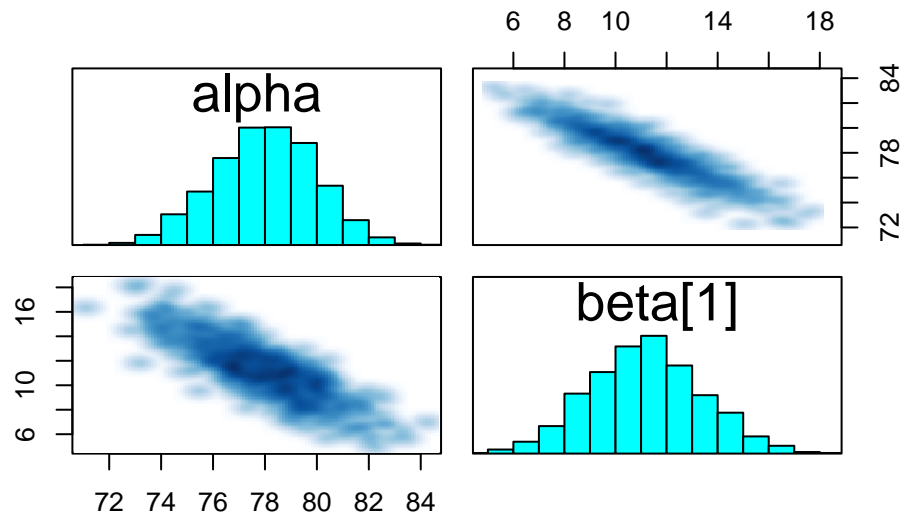
Multiple R-squared: 0.05613, Adjusted R-squared: 0.05394

F-statistic: 25.69 on 1 and 432 DF, p-value: 5.957e-07

We see that the stan estimates alpha and beta are very similar to the intercept and slope from the lm model.

- b) Do a pairs plot to investigate the joint sample distributions of the slope and intercept.
Comment briefly on what you see. Is this potentially a problem?

```
pairs(fit2, pars = c("alpha", "beta[1]"))
```




We see that if we sample high values of α , we are likely to sample low values of β (which is expected when we fit a straight line on the data). This is a potential problem because we will get narrower results when sampling these parameters, which makes the sampling process inefficient.

Plotting results

It might be nice to plot the posterior samples of the estimates for the non-high-school and high-school mothered kids. Here's some code that does this: notice the `beta[condition]` syntax. Also notice I'm using `spread_draws`, because it's easier to calculate the estimated effects in wide format

```
fit2 |>
  spread_draws(alpha, beta[k], sigma) |>
    mutate(nhs = alpha, # no high school is just the intercept
           hs = alpha + beta) |>
  select(nhs, hs) |>
  pivot_longer(nhs:hs, names_to = "education", values_to = "estimated_score") |>
  ggplot(aes(y = education, x = estimated_score)) +
  stat_halfeye() +
  theme_bw() +
```

```
ggtitle("Posterior estimates of scores by education level of mother")
```



Question 4

Add in mother's IQ as a covariate and rerun the model. Please mean center the covariate before putting it into the model. Interpret the coefficient on the (centered) mum's IQ.

```
X <- cbind(as.matrix(kidiq$mom_hs), as.matrix(kidiq$mom_iq - mean(kidiq$mom_iq)))
K <- 2

data <- list(y = y, N = length(y),
             X = X, K = K)
fit3 <- stan(file = here("code/models/kids3.stan"),
             data = data,
             iter = 1000,
             seed = 1)
```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 3.1e-05 seconds

```

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.31 seconds.
Chain 1: Adjust your expectations accordingly!
Chain 1:
Chain 1:
Chain 1: Iteration:   1 / 1000 [  0%] (Warmup)
Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 1:
Chain 1: Elapsed Time: 0.151568 seconds (Warm-up)
Chain 1:                  0.109227 seconds (Sampling)
Chain 1:                  0.260795 seconds (Total)
Chain 1:

```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 2).

```

Chain 2:
Chain 2: Gradient evaluation took 2.6e-05 seconds
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.26 seconds.
Chain 2: Adjust your expectations accordingly!
Chain 2:
Chain 2:
Chain 2: Iteration:   1 / 1000 [  0%] (Warmup)
Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.176378 seconds (Warm-up)

```

Chain 2: 0.104394 seconds (Sampling)
Chain 2: 0.280772 seconds (Total)
Chain 2:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 3).

Chain 3:
Chain 3: Gradient evaluation took 2.6e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.26 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration: 1 / 1000 [0%] (Warmup)
Chain 3: Iteration: 100 / 1000 [10%] (Warmup)
Chain 3: Iteration: 200 / 1000 [20%] (Warmup)
Chain 3: Iteration: 300 / 1000 [30%] (Warmup)
Chain 3: Iteration: 400 / 1000 [40%] (Warmup)
Chain 3: Iteration: 500 / 1000 [50%] (Warmup)
Chain 3: Iteration: 501 / 1000 [50%] (Sampling)
Chain 3: Iteration: 600 / 1000 [60%] (Sampling)
Chain 3: Iteration: 700 / 1000 [70%] (Sampling)
Chain 3: Iteration: 800 / 1000 [80%] (Sampling)
Chain 3: Iteration: 900 / 1000 [90%] (Sampling)
Chain 3: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.17723 seconds (Warm-up)
Chain 3: 0.108693 seconds (Sampling)
Chain 3: 0.285923 seconds (Total)
Chain 3:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 4).

Chain 4:
Chain 4: Gradient evaluation took 2.4e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.24 seconds.
Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration: 1 / 1000 [0%] (Warmup)
Chain 4: Iteration: 100 / 1000 [10%] (Warmup)
Chain 4: Iteration: 200 / 1000 [20%] (Warmup)
Chain 4: Iteration: 300 / 1000 [30%] (Warmup)
Chain 4: Iteration: 400 / 1000 [40%] (Warmup)
Chain 4: Iteration: 500 / 1000 [50%] (Warmup)
Chain 4: Iteration: 501 / 1000 [50%] (Sampling)

```

Chain 4: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 4: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 4: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 4: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 4: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 0.17822 seconds (Warm-up)
Chain 4:           0.103943 seconds (Sampling)
Chain 4:           0.282163 seconds (Total)
Chain 4:

```

```
fit3
```

Inference for Stan model: kids3.

4 chains, each with iter=1000; warmup=500; thin=1;
post-warmup draws per chain=500, total post-warmup draws=2000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
alpha	82.33	0.06	1.89	78.71	81.05	82.37	83.61	86.13
beta[1]	5.67	0.07	2.13	1.48	4.23	5.70	7.12	9.80
beta[2]	0.56	0.00	0.06	0.45	0.53	0.57	0.60	0.68
sigma	18.14	0.02	0.60	17.00	17.73	18.11	18.55	19.35
lp__	-1474.44	0.05	1.42	-1477.94	-1475.13	-1474.08	-1473.40	-1472.67
	n_eff	Rhat						
alpha	970	1.00						
beta[1]	937	1.00						
beta[2]	1280	1.00						
sigma	1447	1.00						
lp__	984	1.01						

Samples were drawn using NUTS(diag_e) at Sat Feb 11 12:23:55 2023.

For each parameter, n_eff is a crude measure of effective sample size,
and Rhat is the potential scale reduction factor on split chains (at
convergence, Rhat=1).

Interpret the coefficient on the (centered) mum's IQ: For every one unit increase in centered mom's IQ, the expected kid's test score increases 0.56, holding all other variables (mom's high school degree status) constant.

Question 5

Confirm the results from Stan agree with lm()

```
kidiq<-kidiq|>
  mutate(mom_iq_center=mom_iq-mean(mom_iq))
summary(lm(kid_score~as.factor(mom_hs)+mom_iq_center,data=kidiq))
```

Call:

```
lm(formula = kid_score ~ as.factor(mom_hs) + mom_iq_center, data = kidiq)
```

Residuals:

Min	1Q	Median	3Q	Max
-52.873	-12.663	2.404	11.356	49.545

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	82.12214	1.94370	42.250	< 2e-16 ***
as.factor(mom_hs)1	5.95012	2.21181	2.690	0.00742 **
mom_iq_center	0.56391	0.06057	9.309	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.14 on 431 degrees of freedom

Multiple R-squared: 0.2141, Adjusted R-squared: 0.2105

F-statistic: 58.72 on 2 and 431 DF, p-value: < 2.2e-16

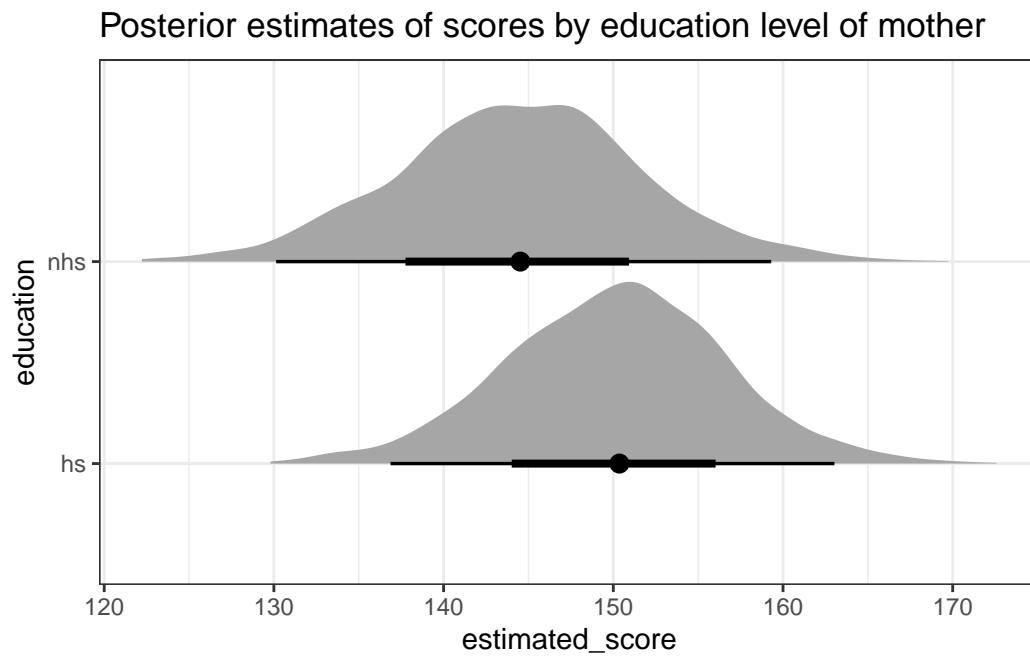
We see that stan's coefficient estimates are similar to the ones from lm model.

Question 6

Plot the posterior estimates of scores by education of mother for mothers who have an IQ of 110.

```
fit3 |>
  spread_draws(alpha, beta[k], sigma) |>
  pivot_wider(names_from = k, values_from = beta, names_glue = "beta{k}") |>
  mutate(nhs=alpha+beta2*110,
         hs=alpha+beta1+beta2*110)|>
  select(nhs, hs) |>
  pivot_longer(nhs:hs, names_to = "education", values_to = "estimated_score") |>
  ggplot(aes(y = education, x = estimated_score)) +
  stat_halfeye() +
```

```
theme_bw() +
ggtitle("Posterior estimates of scores by education level of mother")
```



Question 7

Generate and plot (as a histogram) samples from the posterior predictive distribution for a new kid with a mother who graduated high school and has an IQ of 95.

```
post_samples <- extract(fit3)
alpha <- post_samples[["alpha"]]
beta1 <- post_samples[["beta"]][,1]
beta2 <- post_samples[["beta"]][,2]

lin_pred <- alpha + beta1 + beta2*95
hist(lin_pred)
```