

# Week 5: Bayesian linear regression and introduction to Stan

11/02/23

## Introduction

Today we will be starting off using Stan, looking at the kid's test score data set (available in resources for the [Gelman Hill textbook](#)).

```
library(tidyverse)
library(rstan)
library(tidybayes)
library(here)
```

The data look like this:

```
kidiq <- read_rds(here("data","kidiq.RDS"))
kidiq
```

```
# A tibble: 434 x 4
  kid_score mom_hs mom_iq mom_age
  <int>    <dbl>  <dbl>   <int>
1      65      1  121.     27
2      98      1   89.4     25
3      85      1  115.     27
4      83      1   99.4     25
5     115      1   92.7     27
6      98      0  108.     18
7      69      1  139.     20
8     106      1  125.     23
9     102      1   81.6     24
```

```
10      95      1   95.1      19
# ... with 424 more rows
```

As well as the kid's test scores, we have a binary variable indicating whether or not the mother completed high school, the mother's IQ and age.

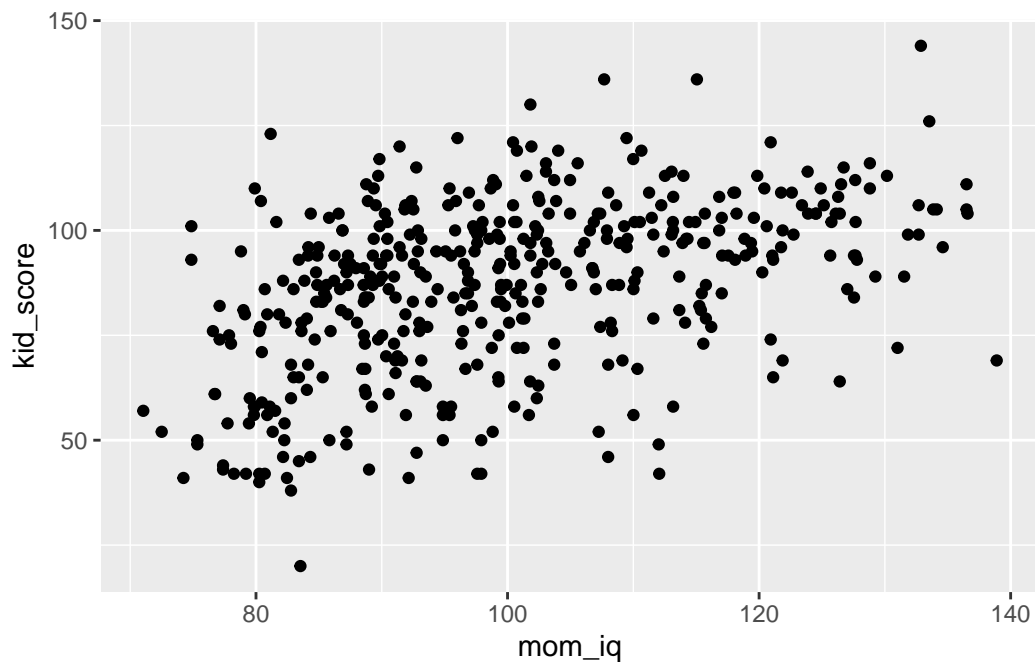
## Descriptives

### Question 1

Use plots or tables to show three interesting observations about the data. Remember:

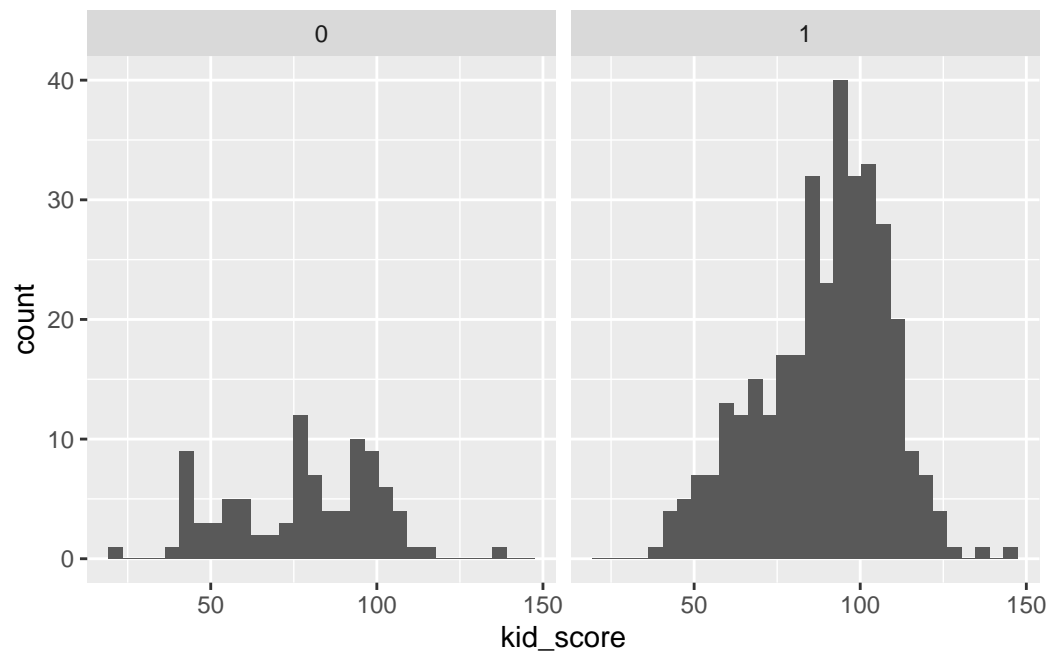
- Explain what your graph/ tables show
- Choose a graph type that's appropriate to the data type

```
ggplot(data=kidiq)+
  geom_point(aes(x=mom_iq, y=kid_score))
```

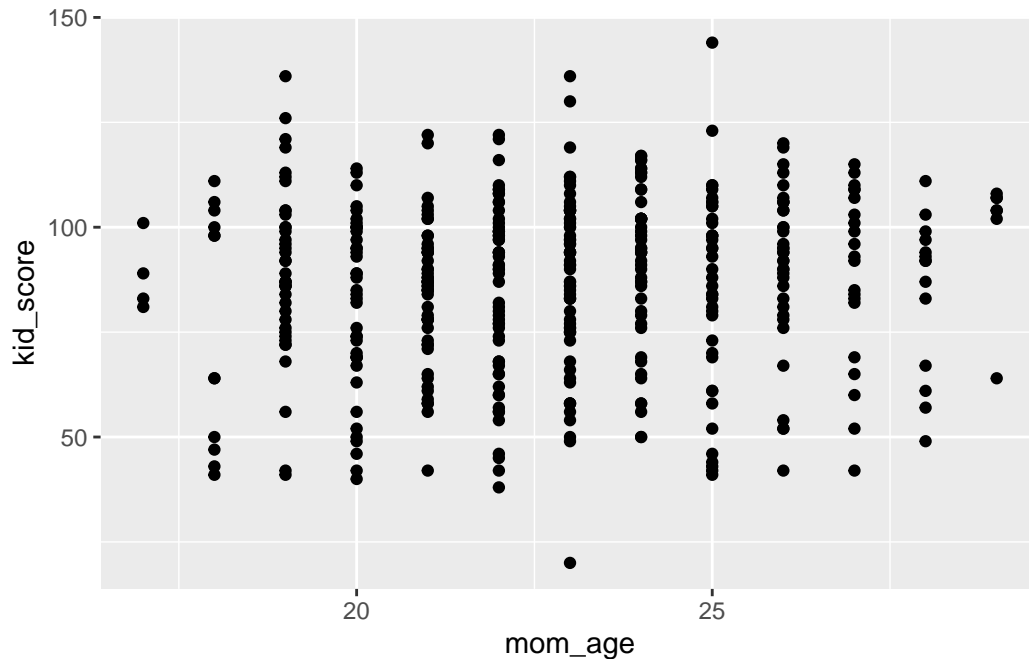


```
ggplot(data=kidiq)+
  geom_histogram(aes(kid_score))+
```

```
facet_grid(~mom_hs)
```



```
ggplot(data=kidiq)+  
  geom_point(aes(x=mom_age, y=kid_score))
```



1. Since both `kid_score` and `mom_iq` are numerical, I first plot a scatter plot of kid's IQ score vs mom's IQ score, and found that kid's IQ score increases as mom's IQ score increases.
2. Since `mom_hs` is binary and `kid_score` is numerical, I then plot histograms of kid's IQ by mom's high school status, and found that for mother with a high school degree, the kid's IQ has larger density on higher IQ scores than for mother without a high school degree.
3. Since both `kid_score` and `mom_age` are numerical, I also plot a scatter plot of kid's IQ score vs mom's age, and found that there is no clear pattern between kid's IQ and mom's age.

## Estimating mean, no covariates

In class we were trying to estimate the mean and standard deviation of the kid's test scores. The `kids2.stan` file contains a Stan model to do this. If you look at it, you will notice the first `data` chunk lists some inputs that we have to define: the outcome variable `y`, number of observations `N`, and the mean and standard deviation of the prior on `mu`. Let's define all these values in a `data` list.

```

y <- kidiq$kid_score
mu0 <- 80
sigma0 <- 10

# named list to input for stan function
data <- list(y = y,
             N = length(y),
             mu0 = mu0,
             sigma0 = sigma0)

```

Now we can run the model:

```

fit <- stan(file = here("code/models/kids2.stan"),
            data = data,
            chains = 3,
            iter = 500,
            seed = 1)

```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 2.9e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.29 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 500 [ 0%] (Warmup)

Chain 1: Iteration: 50 / 500 [ 10%] (Warmup)

Chain 1: Iteration: 100 / 500 [ 20%] (Warmup)

Chain 1: Iteration: 150 / 500 [ 30%] (Warmup)

Chain 1: Iteration: 200 / 500 [ 40%] (Warmup)

Chain 1: Iteration: 250 / 500 [ 50%] (Warmup)

Chain 1: Iteration: 251 / 500 [ 50%] (Sampling)

Chain 1: Iteration: 300 / 500 [ 60%] (Sampling)

Chain 1: Iteration: 350 / 500 [ 70%] (Sampling)

Chain 1: Iteration: 400 / 500 [ 80%] (Sampling)

Chain 1: Iteration: 450 / 500 [ 90%] (Sampling)

Chain 1: Iteration: 500 / 500 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.014204 seconds (Warm-up)

Chain 1: 0.007131 seconds (Sampling)

Chain 1: 0.021335 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 9e-06 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 500 [ 0%] (Warmup)

Chain 2: Iteration: 50 / 500 [ 10%] (Warmup)

Chain 2: Iteration: 100 / 500 [ 20%] (Warmup)

Chain 2: Iteration: 150 / 500 [ 30%] (Warmup)

Chain 2: Iteration: 200 / 500 [ 40%] (Warmup)

Chain 2: Iteration: 250 / 500 [ 50%] (Warmup)

Chain 2: Iteration: 251 / 500 [ 50%] (Sampling)

Chain 2: Iteration: 300 / 500 [ 60%] (Sampling)

Chain 2: Iteration: 350 / 500 [ 70%] (Sampling)

Chain 2: Iteration: 400 / 500 [ 80%] (Sampling)

Chain 2: Iteration: 450 / 500 [ 90%] (Sampling)

Chain 2: Iteration: 500 / 500 [100%] (Sampling)

Chain 2:

Chain 2: Elapsed Time: 0.011218 seconds (Warm-up)

Chain 2: 0.007308 seconds (Sampling)

Chain 2: 0.018526 seconds (Total)

Chain 2:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 3).

Chain 3:

Chain 3: Gradient evaluation took 8e-06 seconds

Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.08 seconds.

Chain 3: Adjust your expectations accordingly!

Chain 3:

Chain 3:

Chain 3: Iteration: 1 / 500 [ 0%] (Warmup)

Chain 3: Iteration: 50 / 500 [ 10%] (Warmup)

Chain 3: Iteration: 100 / 500 [ 20%] (Warmup)

Chain 3: Iteration: 150 / 500 [ 30%] (Warmup)

Chain 3: Iteration: 200 / 500 [ 40%] (Warmup)

Chain 3: Iteration: 250 / 500 [ 50%] (Warmup)

Chain 3: Iteration: 251 / 500 [ 50%] (Sampling)

Chain 3: Iteration: 300 / 500 [ 60%] (Sampling)

Chain 3: Iteration: 350 / 500 [ 70%] (Sampling)

```
Chain 3: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 3: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 3: Iteration: 500 / 500 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.011761 seconds (Warm-up)
Chain 3:           0.007271 seconds (Sampling)
Chain 3:           0.019032 seconds (Total)
Chain 3:
```

Look at the summary

```
fit
```

Inference for Stan model: kids2.

3 chains, each with iter=500; warmup=250; thin=1;

post-warmup draws per chain=250, total post-warmup draws=750.

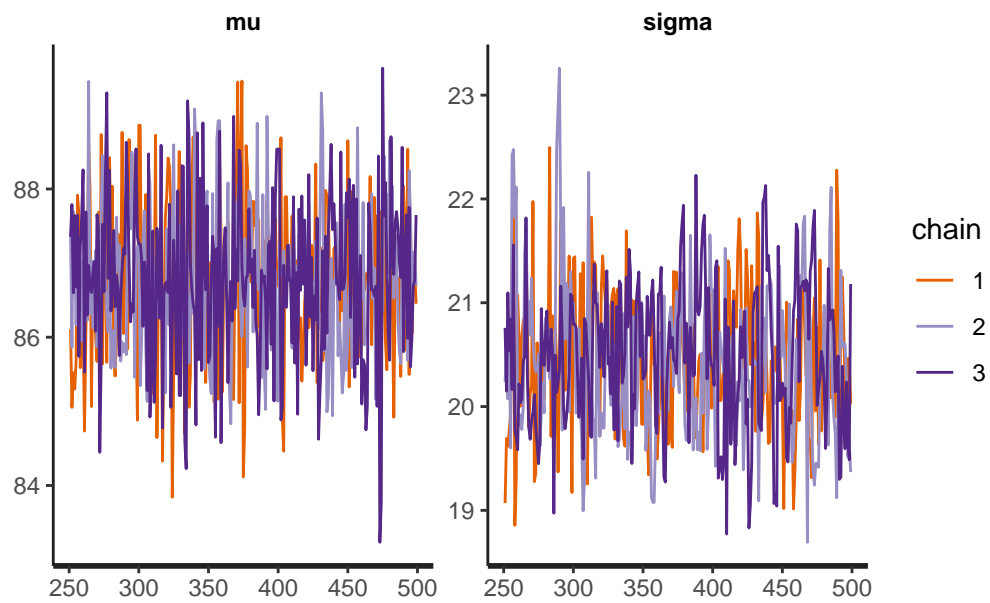
	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
mu	86.79	0.04	1.04	84.81	86.07	86.80	87.48	88.77	743
sigma	20.45	0.04	0.69	19.17	19.97	20.44	20.92	21.87	315
lp__	-1525.83	0.06	1.07	-1528.49	-1526.26	-1525.51	-1525.05	-1524.78	369
Rhat									
mu	1.00								
sigma	1.00								
lp__	1.02								

Samples were drawn using NUTS(diag\_e) at Sat Feb 11 19:47:56 2023.

For each parameter, n\_eff is a crude measure of effective sample size, and Rhat is the potential scale reduction factor on split chains (at convergence, Rhat=1).

Traceplot

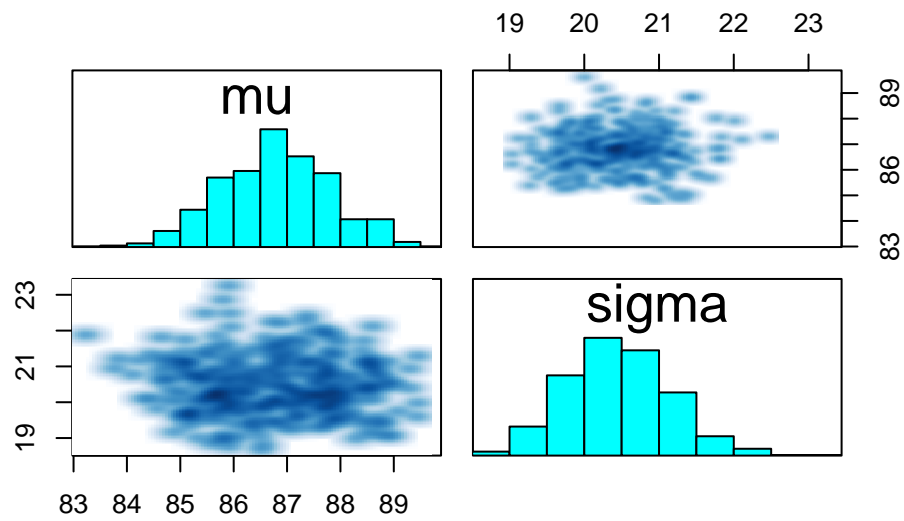
```
traceplot(fit)
```



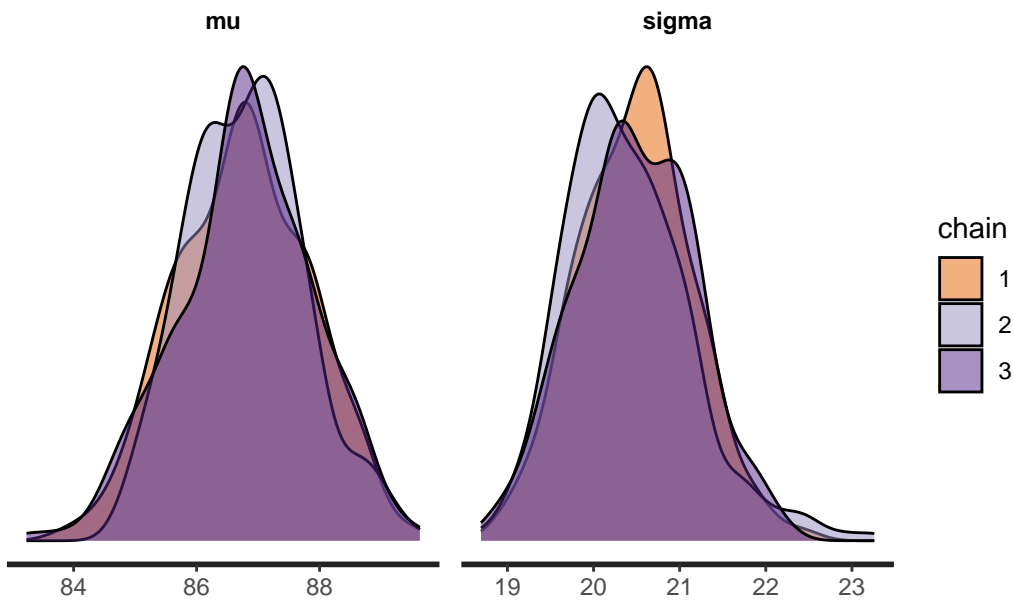
All looks fine.

```
pairs(fit, pars = c("mu", "sigma"))
```





```
stan_dens(fit, separate_chains = TRUE)
```



## Understanding output

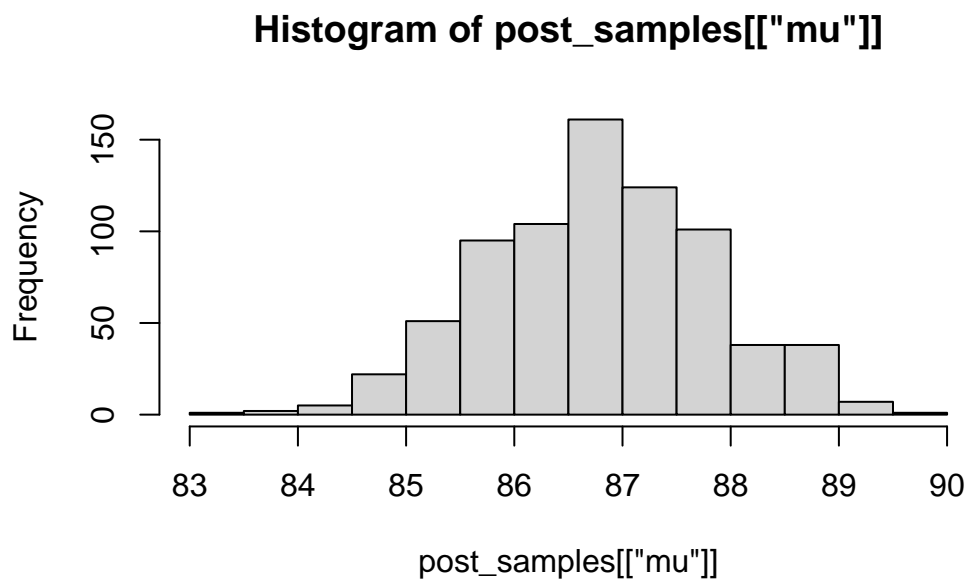
What does the model actually give us? A number of samples from the posteriors. To see this, we can use `extract` to get the samples.

```
post_samples <- extract(fit)
head(post_samples[["mu"]])
```

```
[1] 86.52759 85.49155 86.30489 86.47727 86.79935 86.86390
```

This is a list, and in this case, each element of the list has 4000 samples. E.g. quickly plot a histogram of `mu`

```
hist(post_samples[["mu"]])
```



```
median(post_samples[["mu"]])
```

```
[1] 86.79872
```

```
# 95% bayesian credible interval
quantile(post_samples[["mu"]], 0.025)
```

```
2.5%
84.81158
```

```
quantile(post_samples[["mu"]], 0.975)
```

```
97.5%
88.76504
```

## Plot estimates

There are a bunch of packages, built-in functions that let you plot the estimates from the model, and I encourage you to explore these options (particularly in **bayesplot**, which we will most likely be using later on). I like using the **tidybayes** package, which allows us to easily get the posterior samples in a tidy format (e.g. using `gather_draws` to get in long format). Once we have that, it's easy to just pipe and do ggplots as usual.

Get the posterior samples for mu and sigma in long format:

```
dsamples <- fit |>
  gather_draws(mu, sigma) # gather = long format
dsamples
```

```
# A tibble: 1,500 x 5
# Groups:   .variable [2]
  .chain .iteration .draw .variable .value
  <int>      <int> <int> <chr>      <dbl>
1       1         1     1 mu         86.1
2       1         2     2 mu         85.1
3       1         3     3 mu         85.5
4       1         4     4 mu         85.3
5       1         5     5 mu         85.7
6       1         6     6 mu         87.9
7       1         7     7 mu         87.7
8       1         8     8 mu         85.6
9       1         9     9 mu         86.0
10      1        10    10 mu         86.2
# ... with 1,490 more rows
```

```
# wide format
fit |> spread_draws(mu, sigma)

# A tibble: 750 x 5
  .chain .iteration .draw    mu sigma
  <int>      <int> <int> <dbl> <dbl>
1       1         1     1  86.1  19.1
2       1         2     2  85.1  19.7
3       1         3     3  85.5  19.6
4       1         4     4  85.3  19.8
5       1         5     5  85.7  19.9
6       1         6     6  87.9  21.4
7       1         7     7  87.7  22.1
8       1         8     8  85.6  18.9
9       1         9     9  86.0  19.3
10      1        10    10  86.2  21.1
# ... with 740 more rows
```

```
# quickly calculate the quantiles using
```

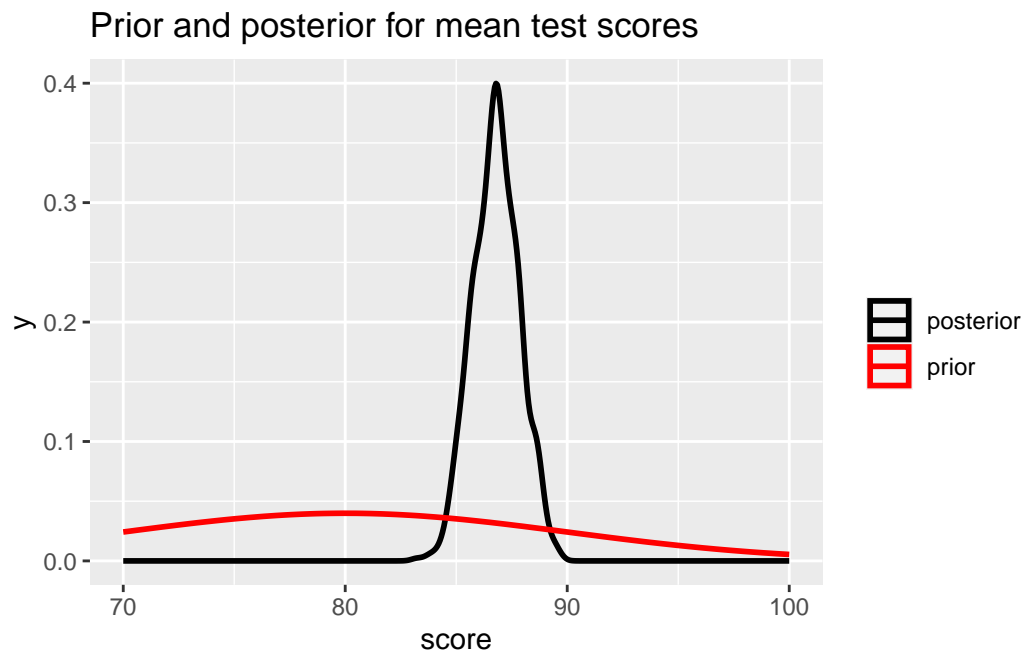
```
dsamples |>
  median_qi(.width = 0.8)
```

```
# A tibble: 2 x 7
  .variable .value .lower .upper .width .point .interval
  <chr>      <dbl> <dbl> <dbl> <dbl> <chr> <chr>
1 mu        86.8  85.5  88.1  0.8 median qi
2 sigma     20.4  19.6  21.3  0.8 median qi
```

Let's plot the density of the posterior samples for mu and add in the prior distribution

```
dsamples |>
  filter(.variable == "mu") |>
  ggplot(aes(.value, color = "posterior")) + geom_density(size = 1) +
  xlim(c(70, 100)) +
  stat_function(fun = dnorm,
    args = list(mean = mu0,
                 sd = sigma0),
    aes(colour = 'prior'), size = 1) +
  scale_color_manual(name = "", values = c("prior" = "red", "posterior" = "black")) +
```

```
ggtitle("Prior and posterior for mean test scores") +
xlab("score")
```



## Question 2

Change the prior to be much more informative (by changing the standard deviation to be 0.1). Rerun the model. Do the estimates change? Plot the prior and posterior densities.

```
y <- kidiq$kid_score
mu0 <- 80
sigma0 <- 0.1

# named list to input for stan function
data <- list(y = y,
             N = length(y),
             mu0 = mu0,
             sigma0 = sigma0)

fit <- stan(file = here("code/models/kids2.stan"),
            data = data,
            chains = 3,
```

```
iter = 500,  
seed = 1)
```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 1.4e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 500 [ 0%] (Warmup)

Chain 1: Iteration: 50 / 500 [ 10%] (Warmup)

Chain 1: Iteration: 100 / 500 [ 20%] (Warmup)

Chain 1: Iteration: 150 / 500 [ 30%] (Warmup)

Chain 1: Iteration: 200 / 500 [ 40%] (Warmup)

Chain 1: Iteration: 250 / 500 [ 50%] (Warmup)

Chain 1: Iteration: 251 / 500 [ 50%] (Sampling)

Chain 1: Iteration: 300 / 500 [ 60%] (Sampling)

Chain 1: Iteration: 350 / 500 [ 70%] (Sampling)

Chain 1: Iteration: 400 / 500 [ 80%] (Sampling)

Chain 1: Iteration: 450 / 500 [ 90%] (Sampling)

Chain 1: Iteration: 500 / 500 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.009571 seconds (Warm-up)

Chain 1: 0.008127 seconds (Sampling)

Chain 1: 0.017698 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 9e-06 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 500 [ 0%] (Warmup)

Chain 2: Iteration: 50 / 500 [ 10%] (Warmup)

Chain 2: Iteration: 100 / 500 [ 20%] (Warmup)

Chain 2: Iteration: 150 / 500 [ 30%] (Warmup)

Chain 2: Iteration: 200 / 500 [ 40%] (Warmup)

Chain 2: Iteration: 250 / 500 [ 50%] (Warmup)

```

Chain 2: Iteration: 251 / 500 [ 50%] (Sampling)
Chain 2: Iteration: 300 / 500 [ 60%] (Sampling)
Chain 2: Iteration: 350 / 500 [ 70%] (Sampling)
Chain 2: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 2: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 2: Iteration: 500 / 500 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.008712 seconds (Warm-up)
Chain 2:           0.007213 seconds (Sampling)
Chain 2:           0.015925 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'kids2' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 9e-06 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.09 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:   1 / 500 [  0%] (Warmup)
Chain 3: Iteration:  50 / 500 [ 10%] (Warmup)
Chain 3: Iteration: 100 / 500 [ 20%] (Warmup)
Chain 3: Iteration: 150 / 500 [ 30%] (Warmup)
Chain 3: Iteration: 200 / 500 [ 40%] (Warmup)
Chain 3: Iteration: 250 / 500 [ 50%] (Warmup)
Chain 3: Iteration: 251 / 500 [ 50%] (Sampling)
Chain 3: Iteration: 300 / 500 [ 60%] (Sampling)
Chain 3: Iteration: 350 / 500 [ 70%] (Sampling)
Chain 3: Iteration: 400 / 500 [ 80%] (Sampling)
Chain 3: Iteration: 450 / 500 [ 90%] (Sampling)
Chain 3: Iteration: 500 / 500 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.009109 seconds (Warm-up)
Chain 3:           0.008079 seconds (Sampling)
Chain 3:           0.017188 seconds (Total)
Chain 3:

```

```
fit
```

Inference for Stan model: kids2.

3 chains, each with iter=500; warmup=250; thin=1;  
post-warmup draws per chain=250, total post-warmup draws=750.

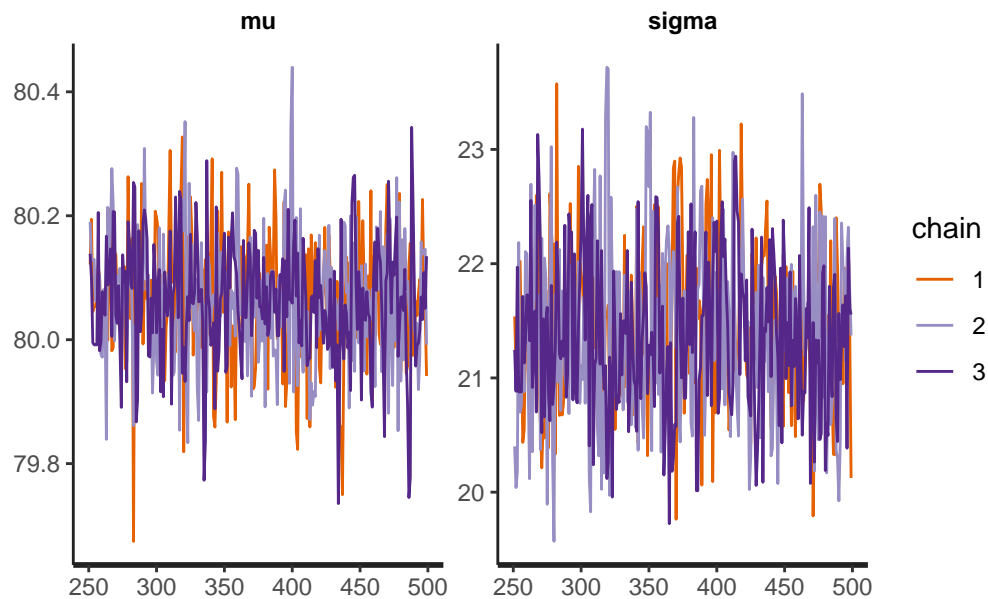
	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff
mu	80.06	0.00	0.10	79.86	80.00	80.07	80.12	80.25	669
sigma	21.42	0.03	0.72	20.09	20.90	21.40	21.90	22.90	618
lp__	-1548.35	0.05	1.03	-1551.09	-1548.67	-1548.04	-1547.67	-1547.40	352

	Rhat
mu	1
sigma	1
lp__	1

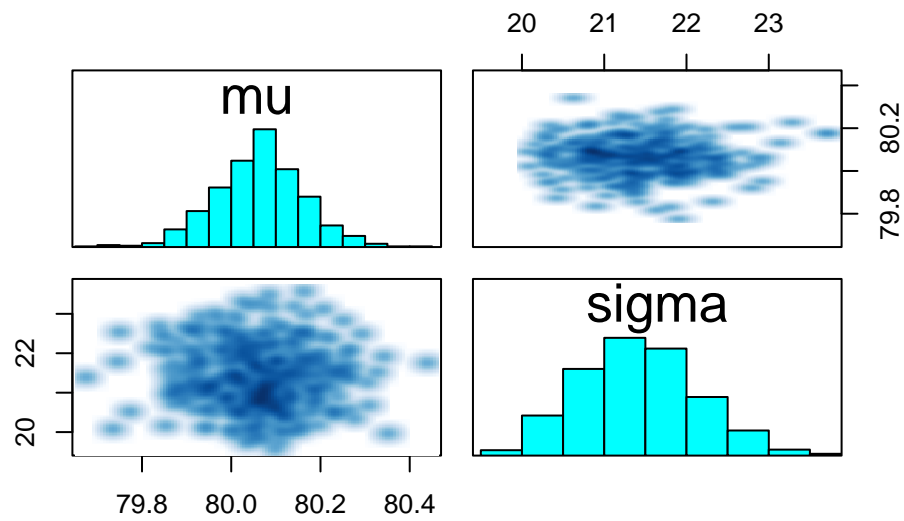
Samples were drawn using NUTS(diag\_e) at Sat Feb 11 19:47:59 2023.  
For each parameter, n\_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).

```
traceplot(fit)
```

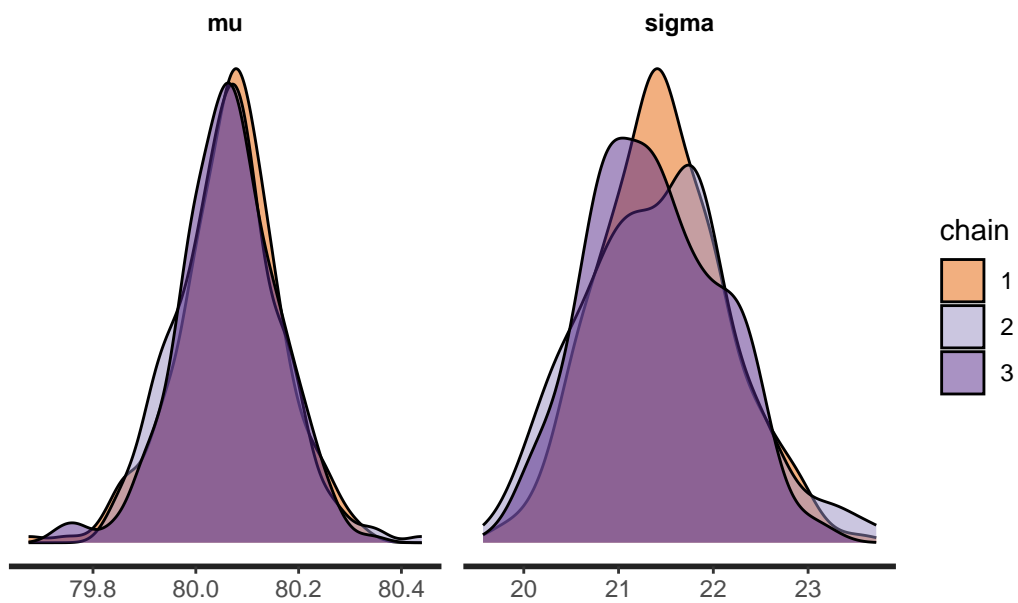


```
pairs(fit, pars = c("mu", "sigma"))
```





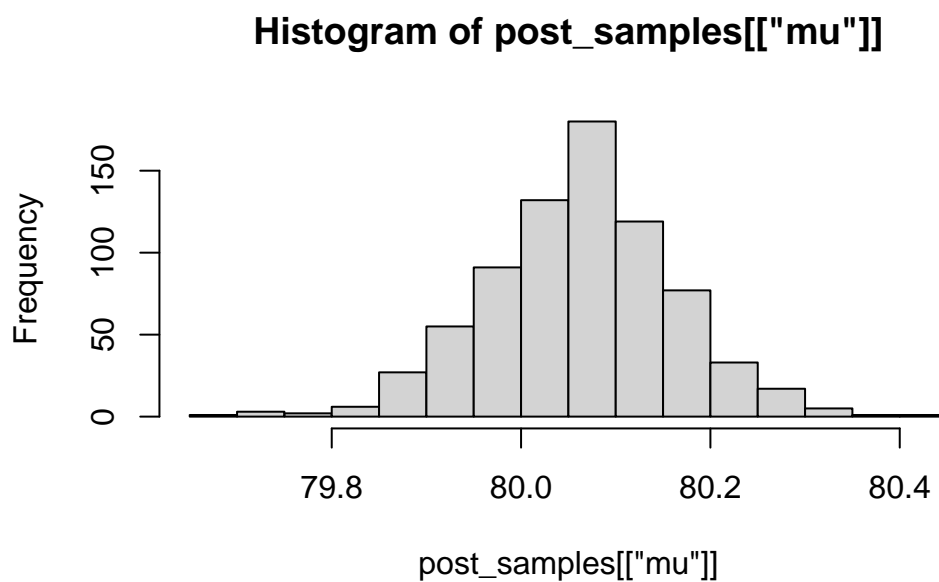
```
stan_dens(fit, separate_chains = TRUE)
```



```
post_samples <- extract(fit)
head(post_samples[["mu"]])
```

```
[1] 80.04107 80.06757 80.09096 80.19419 80.29191 79.67466
```

```
hist(post_samples[["mu"]])
```



```
median(post_samples[["mu"]])
```

```
[1] 80.06706
```

```
# 95% bayesian credible interval
quantile(post_samples[["mu"]], 0.025)
```

```
2.5%
79.86062
```

```
quantile(post_samples[["mu"]], 0.975)
```

```
97.5%  
80.25431
```

```
dsamples <- fit |>  
  gather_draws(mu, sigma) # gather = long format  
dsamples
```

```
# A tibble: 1,500 x 5  
# Groups:   .variable [2]  
  .chain .iteration .draw .variable .value  
  <int>     <int> <int> <chr>     <dbl>  
1       1         1     1 mu       80.1  
2       1         2     2 mu       80.2  
3       1         3     3 mu       80.0  
4       1         4     4 mu       80.1  
5       1         5     5 mu       80.1  
6       1         6     6 mu       80.1  
7       1         7     7 mu       80.1  
8       1         8     8 mu       80.1  
9       1         9     9 mu       80.0  
10      1        10    10 mu       80.1  
# ... with 1,490 more rows
```

```
# wide format  
fit |> spread_draws(mu, sigma)
```

```
# A tibble: 750 x 5  
  .chain .iteration .draw    mu sigma  
  <int>     <int> <int> <dbl> <dbl>  
1       1         1     1  80.1  21.5  
2       1         2     2  80.2  21.1  
3       1         3     3  80.0  21.0  
4       1         4     4  80.1  22.0  
5       1         5     5  80.1  22.0  
6       1         6     6  80.1  21.3  
7       1         7     7  80.1  20.4  
8       1         8     8  80.1  20.5
```

```

  9      1      9      9 80.0 21.4
10      1     10     10 80.1 21.7
# ... with 740 more rows

```

```
# quickly calculate the quantiles using
```

```

dsamples |>
  median_qi(.width = 0.8)

```

```

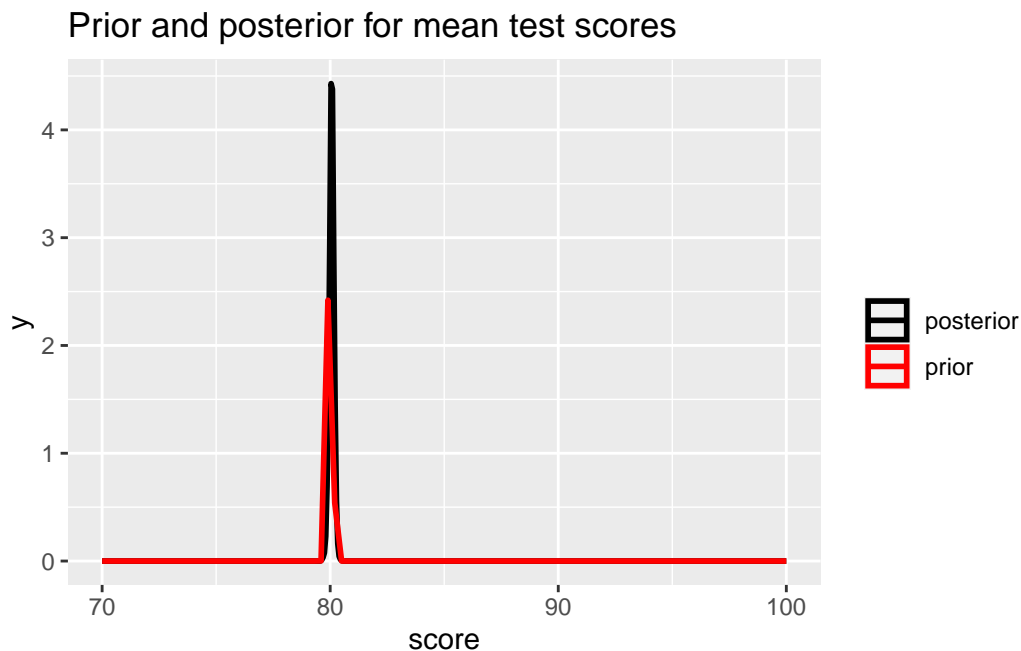
# A tibble: 2 x 7
  .variable .value .lower .upper .width .point .interval
  <chr>      <dbl> <dbl> <dbl> <dbl> <chr> <chr>
1 mu        80.1  79.9  80.2   0.8 median qi
2 sigma     21.4  20.5  22.4   0.8 median qi

```

```

dsamples |>
  filter(.variable == "mu") |>
  ggplot(aes(.value, color = "posterior")) + geom_density(size = 1) +
  xlim(c(70, 100)) +
  stat_function(fun = dnorm,
    args = list(mean = mu0,
                 sd = sigma0),
    aes(colour = 'prior'), size = 1) +
  scale_color_manual(name = "", values = c("prior" = "red", "posterior" = "black")) +
  ggtitle("Prior and posterior for mean test scores") +
  xlab("score")

```



The estimates changed. Previously, the median of  $\mu$  is 86.8 and median of  $\sigma$  is 20.4, but now the median of  $\mu$  is 80.1 and median of  $\sigma$  is 21.4.

## Adding covariates

Now let's see how kid's test scores are related to mother's education. We want to run the simple linear regression

$$Score = \alpha + \beta X$$

where  $X = 1$  if the mother finished high school and zero otherwise.

`kid3.stan` has the stan model to do this. Notice now we have some inputs related to the design matrix  $X$  and the number of covariates (in this case, it's just 1).

Let's get the data we need and run the model.

```
X <- as.matrix(kidiq$mom_hs, ncol = 1) # force this to be a matrix
K <- 1

data <- list(y = y, N = length(y),
             X = X, K = K)
```

```
fit2 <- stan(file = here("code/models/kids3.stan"),
             data = data,
             iter = 1000,
             seed = 1)
```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 7.5e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.75 seconds.

Chain 1: Adjust your expectations accordingly!

Chain 1:

Chain 1:

Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)

Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)

Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)

Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)

Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)

Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)

Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)

Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)

Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)

Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)

Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)

Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)

Chain 1:

Chain 1: Elapsed Time: 0.183939 seconds (Warm-up)

Chain 1: 0.134833 seconds (Sampling)

Chain 1: 0.318772 seconds (Total)

Chain 1:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 2).

Chain 2:

Chain 2: Gradient evaluation took 3.7e-05 seconds

Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.37 seconds.

Chain 2: Adjust your expectations accordingly!

Chain 2:

Chain 2:

Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)

Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)

Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)

Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)

```

Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 2:
Chain 2: Elapsed Time: 0.188002 seconds (Warm-up)
Chain 2:                0.120174 seconds (Sampling)
Chain 2:                0.308176 seconds (Total)
Chain 2:

```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 3).

```

Chain 3:
Chain 3: Gradient evaluation took 3.3e-05 seconds
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.33 seconds.
Chain 3: Adjust your expectations accordingly!
Chain 3:
Chain 3:
Chain 3: Iteration:   1 / 1000 [  0%] (Warmup)
Chain 3: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 3: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 3: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 3: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 3: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 3: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 3: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 3: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 3: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 3: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 3: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 3:
Chain 3: Elapsed Time: 0.231008 seconds (Warm-up)
Chain 3:                0.13619 seconds (Sampling)
Chain 3:                0.367198 seconds (Total)
Chain 3:

```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 4).

```

Chain 4:
Chain 4: Gradient evaluation took 3.8e-05 seconds
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.38 seconds.

```

```

Chain 4: Adjust your expectations accordingly!
Chain 4:
Chain 4:
Chain 4: Iteration: 1 / 1000 [ 0%] (Warmup)
Chain 4: Iteration: 100 / 1000 [ 10%] (Warmup)
Chain 4: Iteration: 200 / 1000 [ 20%] (Warmup)
Chain 4: Iteration: 300 / 1000 [ 30%] (Warmup)
Chain 4: Iteration: 400 / 1000 [ 40%] (Warmup)
Chain 4: Iteration: 500 / 1000 [ 50%] (Warmup)
Chain 4: Iteration: 501 / 1000 [ 50%] (Sampling)
Chain 4: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 4: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 4: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 4: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 4: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 0.182679 seconds (Warm-up)
Chain 4: 0.127646 seconds (Sampling)
Chain 4: 0.310325 seconds (Total)
Chain 4:

```

### Question 3

- a) Confirm that the estimates of the intercept and slope are comparable to results from `lm()`

```
fit2
```

Inference for Stan model: kids3.

4 chains, each with iter=1000; warmup=500; thin=1;  
post-warmup draws per chain=500, total post-warmup draws=2000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
alpha	78.00	0.07	1.94	74.12	76.74	78.06	79.33	81.67
beta[1]	11.18	0.08	2.19	7.00	9.66	11.17	12.60	15.56
sigma	19.83	0.02	0.66	18.59	19.39	19.81	20.25	21.23
lp__	-1514.33	0.04	1.21	-1517.35	-1514.93	-1514.03	-1513.44	-1512.98
	n_eff	Rhat						
alpha	804	1.01						
beta[1]	846	1.01						
sigma	1042	1.00						



```
lp__      758 1.00
```

Samples were drawn using NUTS(diag\_e) at Sat Feb 11 19:48:34 2023.  
For each parameter, n\_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).

```
model<-lm(kid_score~as.factor(mom_hs), data=kidiq)
summary(model)
```

Call:

```
lm(formula = kid_score ~ as.factor(mom_hs), data = kidiq)
```

Residuals:

Min	1Q	Median	3Q	Max
-57.55	-13.32	2.68	14.68	58.45

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	77.548	2.059	37.670	< 2e-16 ***
as.factor(mom_hs)1	11.771	2.322	5.069	5.96e-07 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 19.85 on 432 degrees of freedom

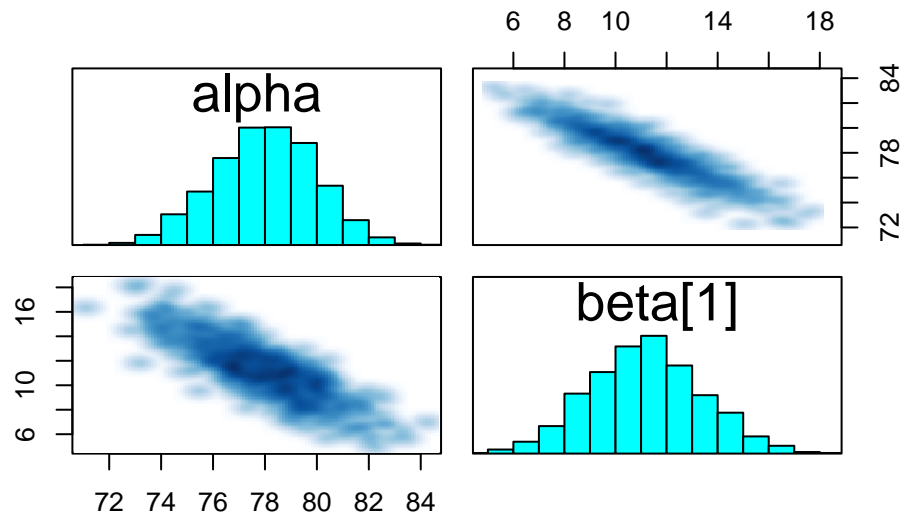
Multiple R-squared: 0.05613, Adjusted R-squared: 0.05394

F-statistic: 25.69 on 1 and 432 DF, p-value: 5.957e-07

We see that the stan estimates alpha and beta are very similar to the intercept and slope from the lm model.

- b) Do a pairs plot to investigate the joint sample distributions of the slope and intercept.  
Comment briefly on what you see. Is this potentially a problem?

```
pairs(fit2, pars = c("alpha", "beta[1]"))
```



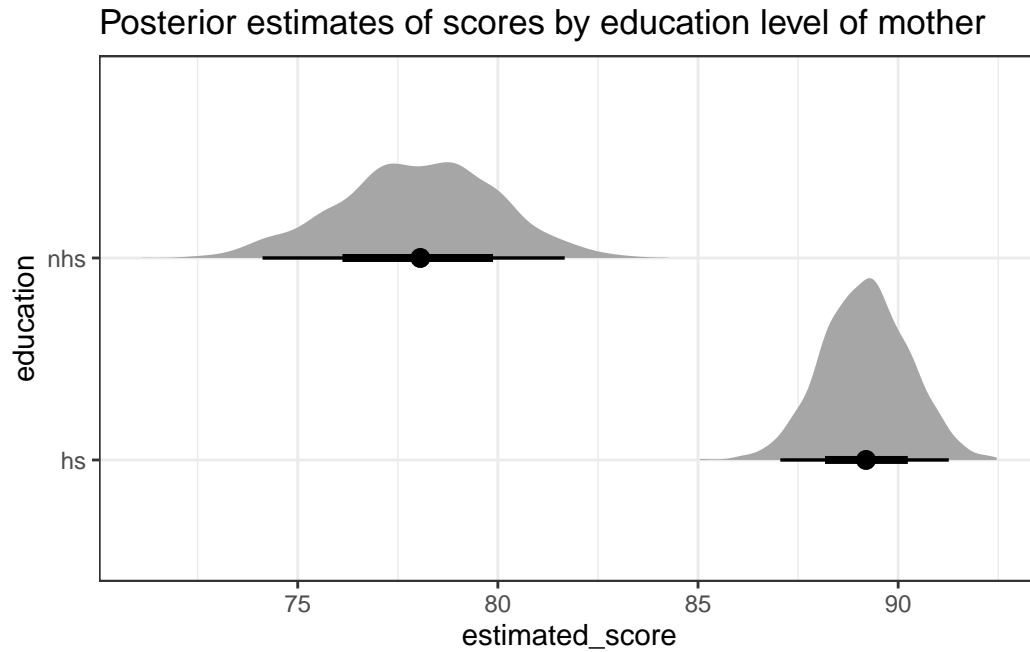
We see that if we sample high values of  $\alpha$ , we are likely to sample low values of  $\beta$  (which is expected when we fit a straight line on the data). This is a potential problem because we will get narrower results when sampling these parameters, which makes the sampling process inefficient.

## Plotting results

It might be nice to plot the posterior samples of the estimates for the non-high-school and high-school mothered kids. Here's some code that does this: notice the `beta[condition]` syntax. Also notice I'm using `spread_draws`, because it's easier to calculate the estimated effects in wide format

```
fit2 |>
  spread_draws(alpha, beta[k], sigma) |>
    mutate(nhs = alpha, # no high school is just the intercept
           hs = alpha + beta) |>
  select(nhs, hs) |>
  pivot_longer(nhs:hs, names_to = "education", values_to = "estimated_score") |>
  ggplot(aes(y = education, x = estimated_score)) +
  stat_halfeye() +
  theme_bw() +
```

```
ggtitle("Posterior estimates of scores by education level of mother")
```



#### Question 4

Add in mother's IQ as a covariate and rerun the model. Please mean center the covariate before putting it into the model. Interpret the coefficient on the (centered) mum's IQ.

```
X <- cbind(as.matrix(kidiq$mom_hs), as.matrix(kidiq$mom_iq - mean(kidiq$mom_iq)))
K <- 2

data <- list(y = y, N = length(y),
             X = X, K = K)
fit3 <- stan(file = here("code/models/kids3.stan"),
             data = data,
             iter = 1000,
             seed = 1)
```

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 1).

Chain 1:

Chain 1: Gradient evaluation took 3.8e-05 seconds

Chain 1: 1000 transitions using 10 leapfrog steps per transition would take 0.38 seconds.  
Chain 1: Adjust your expectations accordingly!  
Chain 1:  
Chain 1:  
Chain 1: Iteration: 1 / 1000 [ 0%] (Warmup)  
Chain 1: Iteration: 100 / 1000 [ 10%] (Warmup)  
Chain 1: Iteration: 200 / 1000 [ 20%] (Warmup)  
Chain 1: Iteration: 300 / 1000 [ 30%] (Warmup)  
Chain 1: Iteration: 400 / 1000 [ 40%] (Warmup)  
Chain 1: Iteration: 500 / 1000 [ 50%] (Warmup)  
Chain 1: Iteration: 501 / 1000 [ 50%] (Sampling)  
Chain 1: Iteration: 600 / 1000 [ 60%] (Sampling)  
Chain 1: Iteration: 700 / 1000 [ 70%] (Sampling)  
Chain 1: Iteration: 800 / 1000 [ 80%] (Sampling)  
Chain 1: Iteration: 900 / 1000 [ 90%] (Sampling)  
Chain 1: Iteration: 1000 / 1000 [100%] (Sampling)  
Chain 1:  
Chain 1: Elapsed Time: 0.214734 seconds (Warm-up)  
Chain 1: 0.160744 seconds (Sampling)  
Chain 1: 0.375478 seconds (Total)  
Chain 1:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 2).

Chain 2:  
Chain 2: Gradient evaluation took 3.2e-05 seconds  
Chain 2: 1000 transitions using 10 leapfrog steps per transition would take 0.32 seconds.  
Chain 2: Adjust your expectations accordingly!  
Chain 2:  
Chain 2:  
Chain 2: Iteration: 1 / 1000 [ 0%] (Warmup)  
Chain 2: Iteration: 100 / 1000 [ 10%] (Warmup)  
Chain 2: Iteration: 200 / 1000 [ 20%] (Warmup)  
Chain 2: Iteration: 300 / 1000 [ 30%] (Warmup)  
Chain 2: Iteration: 400 / 1000 [ 40%] (Warmup)  
Chain 2: Iteration: 500 / 1000 [ 50%] (Warmup)  
Chain 2: Iteration: 501 / 1000 [ 50%] (Sampling)  
Chain 2: Iteration: 600 / 1000 [ 60%] (Sampling)  
Chain 2: Iteration: 700 / 1000 [ 70%] (Sampling)  
Chain 2: Iteration: 800 / 1000 [ 80%] (Sampling)  
Chain 2: Iteration: 900 / 1000 [ 90%] (Sampling)  
Chain 2: Iteration: 1000 / 1000 [100%] (Sampling)  
Chain 2:  
Chain 2: Elapsed Time: 0.259357 seconds (Warm-up)

Chain 2: 0.148798 seconds (Sampling)  
Chain 2: 0.408155 seconds (Total)  
Chain 2:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 3).

Chain 3:  
Chain 3: Gradient evaluation took 3.2e-05 seconds  
Chain 3: 1000 transitions using 10 leapfrog steps per transition would take 0.32 seconds.  
Chain 3: Adjust your expectations accordingly!  
Chain 3:  
Chain 3:  
Chain 3: Iteration: 1 / 1000 [ 0%] (Warmup)  
Chain 3: Iteration: 100 / 1000 [ 10%] (Warmup)  
Chain 3: Iteration: 200 / 1000 [ 20%] (Warmup)  
Chain 3: Iteration: 300 / 1000 [ 30%] (Warmup)  
Chain 3: Iteration: 400 / 1000 [ 40%] (Warmup)  
Chain 3: Iteration: 500 / 1000 [ 50%] (Warmup)  
Chain 3: Iteration: 501 / 1000 [ 50%] (Sampling)  
Chain 3: Iteration: 600 / 1000 [ 60%] (Sampling)  
Chain 3: Iteration: 700 / 1000 [ 70%] (Sampling)  
Chain 3: Iteration: 800 / 1000 [ 80%] (Sampling)  
Chain 3: Iteration: 900 / 1000 [ 90%] (Sampling)  
Chain 3: Iteration: 1000 / 1000 [100%] (Sampling)  
Chain 3:  
Chain 3: Elapsed Time: 0.252556 seconds (Warm-up)  
Chain 3: 0.156238 seconds (Sampling)  
Chain 3: 0.408794 seconds (Total)  
Chain 3:

SAMPLING FOR MODEL 'kids3' NOW (CHAIN 4).

Chain 4:  
Chain 4: Gradient evaluation took 3.5e-05 seconds  
Chain 4: 1000 transitions using 10 leapfrog steps per transition would take 0.35 seconds.  
Chain 4: Adjust your expectations accordingly!  
Chain 4:  
Chain 4:  
Chain 4: Iteration: 1 / 1000 [ 0%] (Warmup)  
Chain 4: Iteration: 100 / 1000 [ 10%] (Warmup)  
Chain 4: Iteration: 200 / 1000 [ 20%] (Warmup)  
Chain 4: Iteration: 300 / 1000 [ 30%] (Warmup)  
Chain 4: Iteration: 400 / 1000 [ 40%] (Warmup)  
Chain 4: Iteration: 500 / 1000 [ 50%] (Warmup)  
Chain 4: Iteration: 501 / 1000 [ 50%] (Sampling)

```

Chain 4: Iteration: 600 / 1000 [ 60%] (Sampling)
Chain 4: Iteration: 700 / 1000 [ 70%] (Sampling)
Chain 4: Iteration: 800 / 1000 [ 80%] (Sampling)
Chain 4: Iteration: 900 / 1000 [ 90%] (Sampling)
Chain 4: Iteration: 1000 / 1000 [100%] (Sampling)
Chain 4:
Chain 4: Elapsed Time: 0.253768 seconds (Warm-up)
Chain 4:           0.149135 seconds (Sampling)
Chain 4:           0.402903 seconds (Total)
Chain 4:

```

```
fit3
```

Inference for Stan model: kids3.

4 chains, each with iter=1000; warmup=500; thin=1;  
post-warmup draws per chain=500, total post-warmup draws=2000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%
alpha	82.33	0.06	1.89	78.71	81.05	82.37	83.61	86.13
beta[1]	5.67	0.07	2.13	1.48	4.23	5.70	7.12	9.80
beta[2]	0.56	0.00	0.06	0.45	0.53	0.57	0.60	0.68
sigma	18.14	0.02	0.60	17.00	17.73	18.11	18.55	19.35
lp__	-1474.44	0.05	1.42	-1477.94	-1475.13	-1474.08	-1473.40	-1472.67
	n_eff	Rhat						
alpha	970	1.00						
beta[1]	937	1.00						
beta[2]	1280	1.00						
sigma	1447	1.00						
lp__	984	1.01						

Samples were drawn using NUTS(diag\_e) at Sat Feb 11 19:48:36 2023.

For each parameter, n\_eff is a crude measure of effective sample size,  
and Rhat is the potential scale reduction factor on split chains (at  
convergence, Rhat=1).

Interpret the coefficient on the (centered) mum's IQ: For every one unit increase in centered mom's IQ, the expected kid's test score increases 0.56, holding all other variables (mom's high school degree status) constant.

## Question 5

Confirm the results from Stan agree with lm()

```
kidiq<-kidiq|>
  mutate(mom_iq_center=mom_iq-mean(mom_iq))
summary(lm(kid_score~as.factor(mom_hs)+mom_iq_center,data=kidiq))
```

Call:

```
lm(formula = kid_score ~ as.factor(mom_hs) + mom_iq_center, data = kidiq)
```

Residuals:

Min	1Q	Median	3Q	Max
-52.873	-12.663	2.404	11.356	49.545

Coefficients:

	Estimate	Std. Error	t value	Pr(> t )
(Intercept)	82.12214	1.94370	42.250	< 2e-16 ***
as.factor(mom_hs)1	5.95012	2.21181	2.690	0.00742 **
mom_iq_center	0.56391	0.06057	9.309	< 2e-16 ***

---

Signif. codes: 0 '\*\*\*' 0.001 '\*\*' 0.01 '\*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 18.14 on 431 degrees of freedom

Multiple R-squared: 0.2141, Adjusted R-squared: 0.2105

F-statistic: 58.72 on 2 and 431 DF, p-value: < 2.2e-16

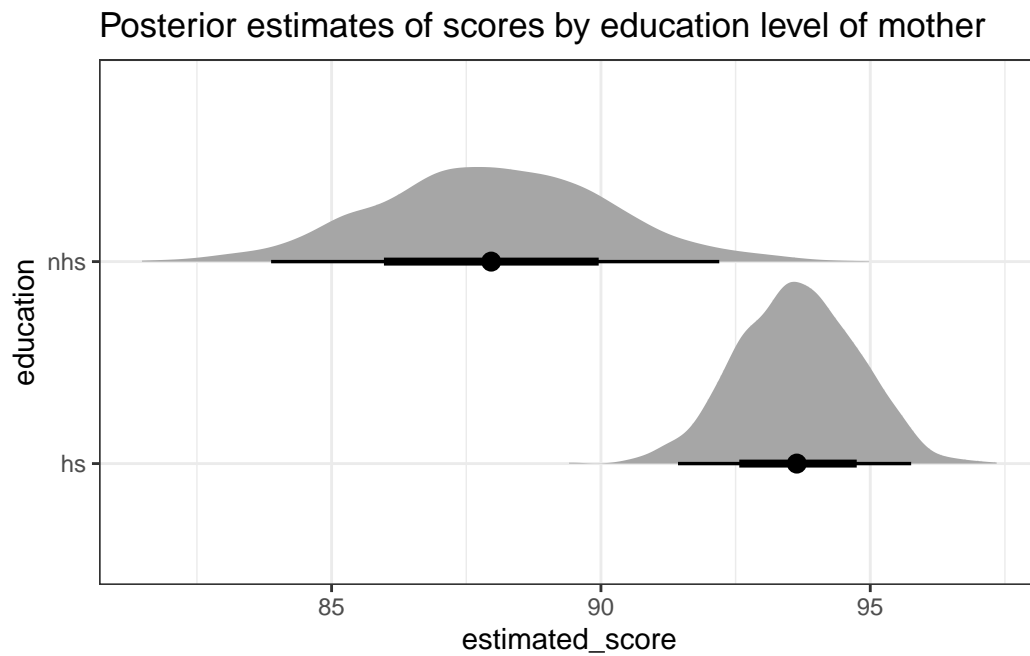
We see that stan's coefficient estimates are similar to the ones from lm model.

## Question 6

Plot the posterior estimates of scores by education of mother for mothers who have an IQ of 110.

```
fit3 |>
  spread_draws(alpha, beta[k], sigma) |>
  pivot_wider(names_from = k, values_from = beta, names_glue = "beta{k}") |>
  mutate(nhs=alpha+beta2*(110-mean(kidiq$mom_iq)),
         hs=alpha+beta1+beta2*(110-mean(kidiq$mom_iq)))|>
  select(nhs, hs) |>
  pivot_longer(nhs:hs, names_to = "education", values_to = "estimated_score") |>
  ggplot(aes(y = education, x = estimated_score)) +
  stat_halfeye() +
```

```
theme_bw() +
ggtitle("Posterior estimates of scores by education level of mother")
```



### Question 7

Generate and plot (as a histogram) samples from the posterior predictive distribution for a new kid with a mother who graduated high school and has an IQ of 95.

```
post_samples <- extract(fit3)
alpha <- post_samples[["alpha"]]
beta1 <- post_samples[["beta"]][,1]
beta2 <- post_samples[["beta"]][,2]

lin_pred <- alpha + beta1 + beta2*(95-mean(kidiq$mom_iq))
hist(lin_pred)
```



